

# Trusted Execution Environments (TEEs)

Shweta Shinde

# Motivating Example



**Input**



**Private Key**



Bank's Web Server

**Signature**



Sensitive

# What can go wrong: Attack Surface

- Signature Verification
  - bad crypto (design / implementation)
- Webserver <-> client
  - broken communication channel
  - MITM, broken protocol (design/implementation)
- Web server
  - software bugs
- Browser
  - software bugs
- Other avenues of attack?

# Google patches two new zero-day flaws in Chrome

The last three weeks have seen a bumper crop of patches for zero-day bugs across software from Google, Apple and Microsoft



Amer Owaida

12 Nov 2020 - 05:25PM



Ben Hawkes  
@benhawkes



In addition to last week's Chrome/freetype 0day (CVE-2020-15999), Project Zero also detected and reported the Windows kernel bug (CVE-2020-17087) that was used for a sandbox escape. The technical details of CVE-2020-17087 are now available here: [bugs.chromium.org/p/project-zero...](https://bugs.chromium.org/p/project-zero/)

5:00 PM · Oct 30, 2020



363 183 people are Tweeting about this

Windows kernel bug (CVE-2020-17087) that was used for a sandbox escape. The technical details of CVE-2020-17087

Apple Safari with a macOS kernel escalation of privilege.

# Team IDs Real-world Vulnerabilities In Popular Browser During Premier Hackathon



Monday, April 6, 2020

## WELCOME TO PWN2OWN 2020 - THE SCHEDULE AND LIVE RESULTS

March 18, 2020 | Dustin Childs

Welcome to Pwn2Own 2020! This year's contest will be the first where all attempts occur remotely. We have contestants from around the world ready to demonstrate some amazing research, and we have ZDI researchers ready to run and verify their attempts. Our returning partner Microsoft and sponsor VMware will also be online with us, as well as other affected vendors. We might not be in the same room, but we're virtually together and ready for a fantastic contest.

The schedule for today is posted below, and we'll be updating this blog throughout the day with results and updates. We'll post a full summary of today's events (including videos of the attempts) tomorrow morning.

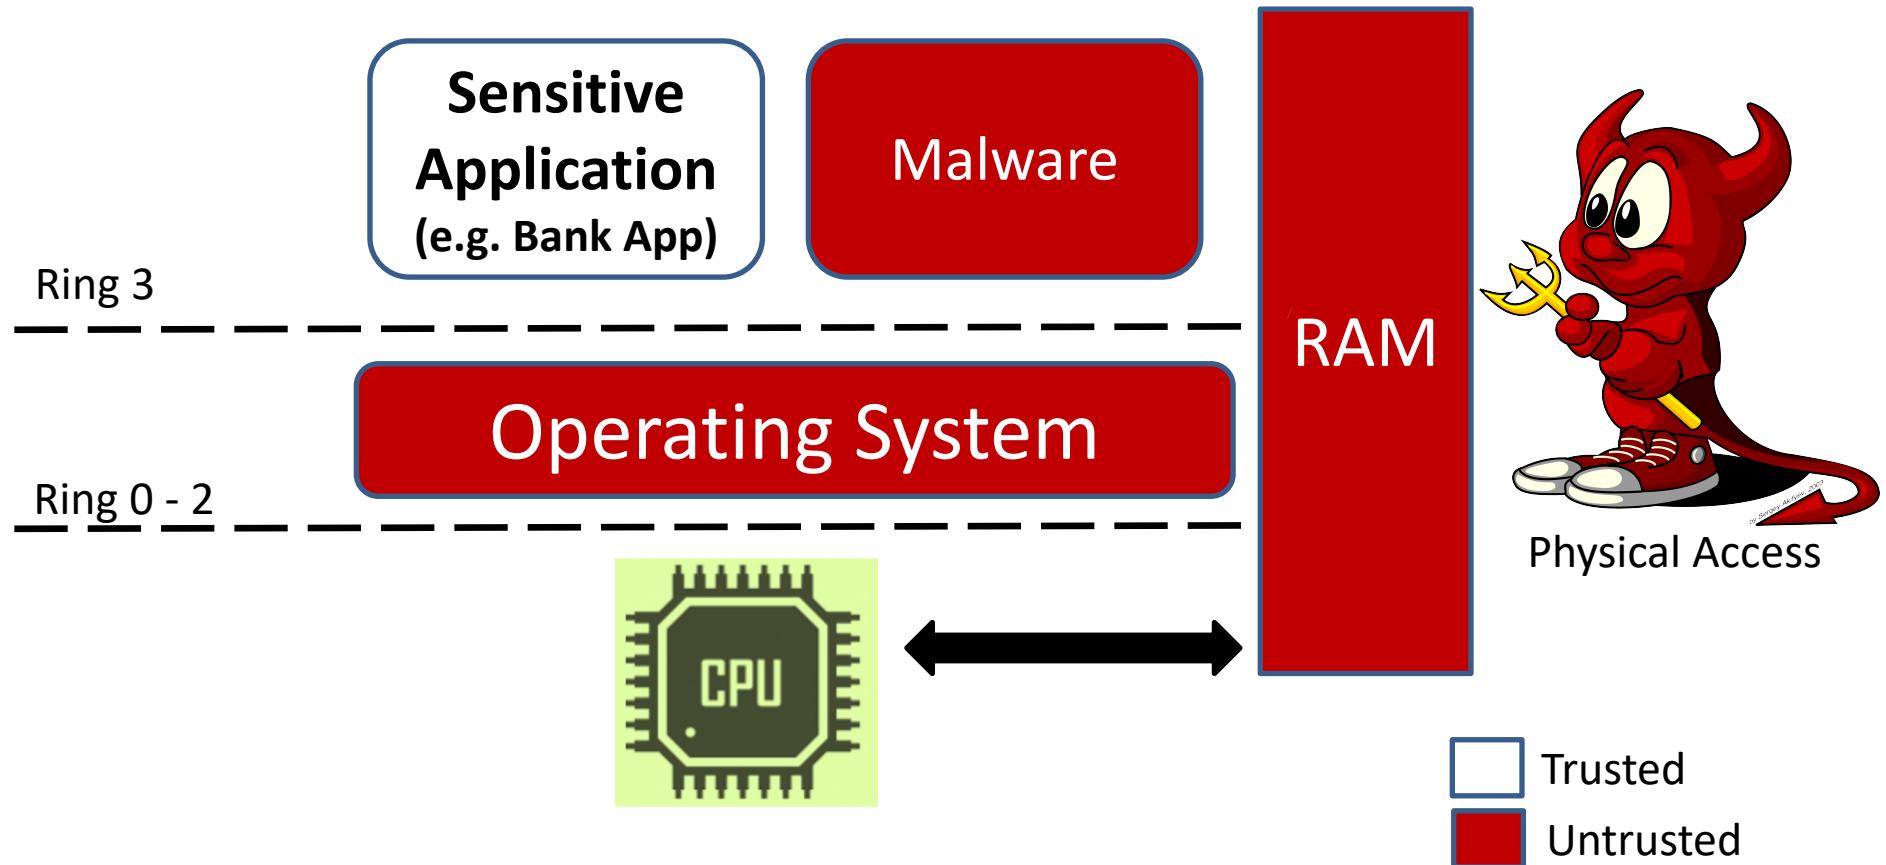
As always, we started the contest with a random drawing to determine the order of attempts. We have four attempts scheduled for today and four queued up for tomorrow. The full schedule for Day One is below (all times Pacific [UTC -7:00]). We will update this schedule with results as they become available.

### Day One – March 18, 2020

1000 – The Georgia Tech Systems Software & Security Lab (@SSLab\_Gatech) team of Yong Hwi Jin (@jinmo123), Jungwon Lim (@setuid0x0\_), and Insu Yun (@insu\_yun\_en) targeting Apple Safari with a macOS kernel escalation of privilege.

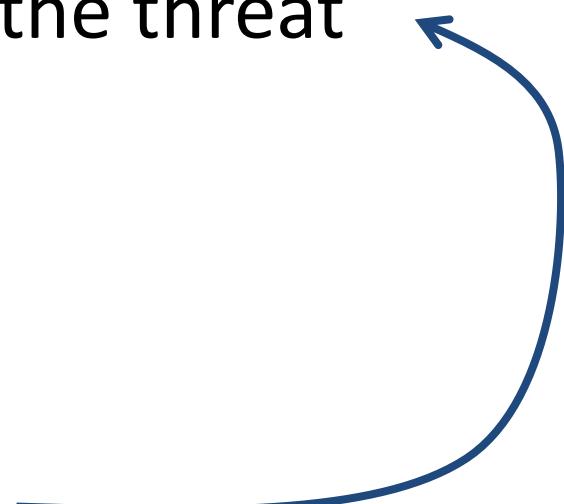
# A Much Larger Attack Surface

- Threats to your bank application



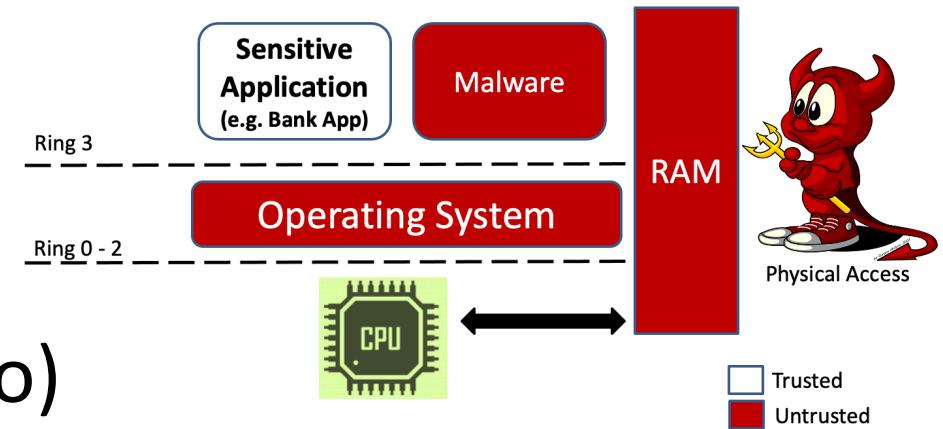
# This Class: the good, the bad, the ugly

- Look at a new threat model and its ramifications
- Learn how to build systematic defenses against the threat
- Understand the gaps in our defenses
- Leverage the gaps to circumvent the defenses



# Possible Defenses

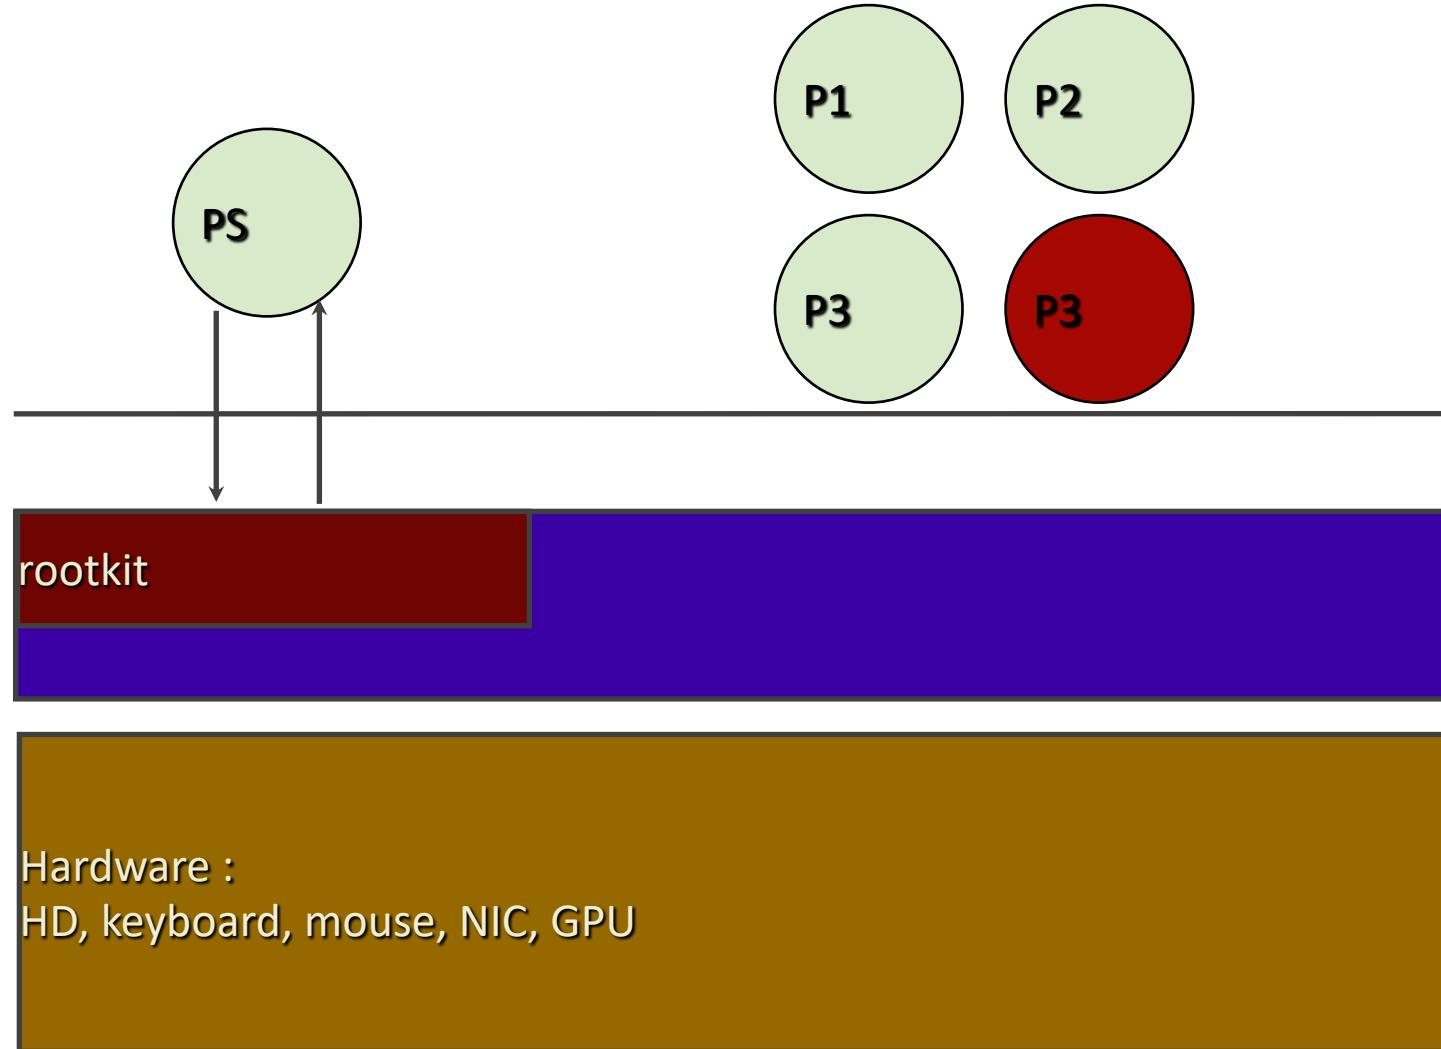
- Binary analysis (e.g., Fuzzing)
- Binary hardening (e.g., CFI, SFI)
- Memory safe languages (e.g., Rust, Go)
- Verification (e.g., seL4)
- Encrypted computation (e.g., homomorphic encryption)
  - How to protect the private keys?
- Virtualization (e.g., XEN)
  - Assumes hypervisor is uncompromised



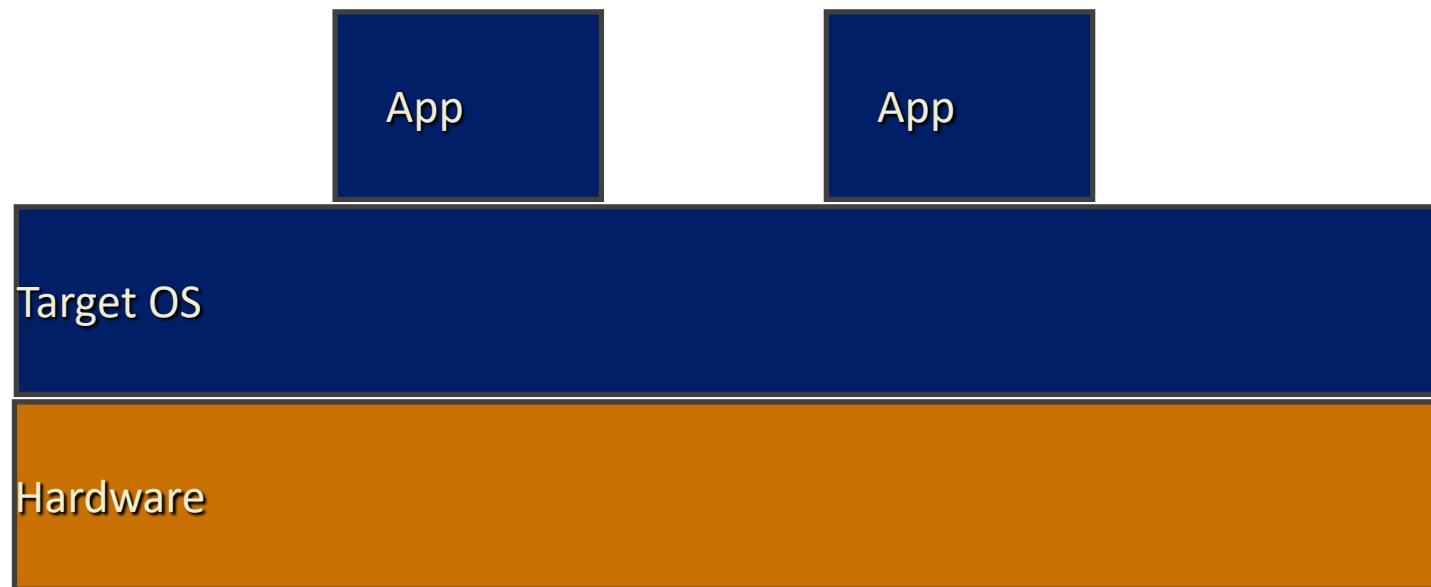
These approaches are not scalable and/or complete

How can one compromise the confidentiality & integrity of sensitive computation?

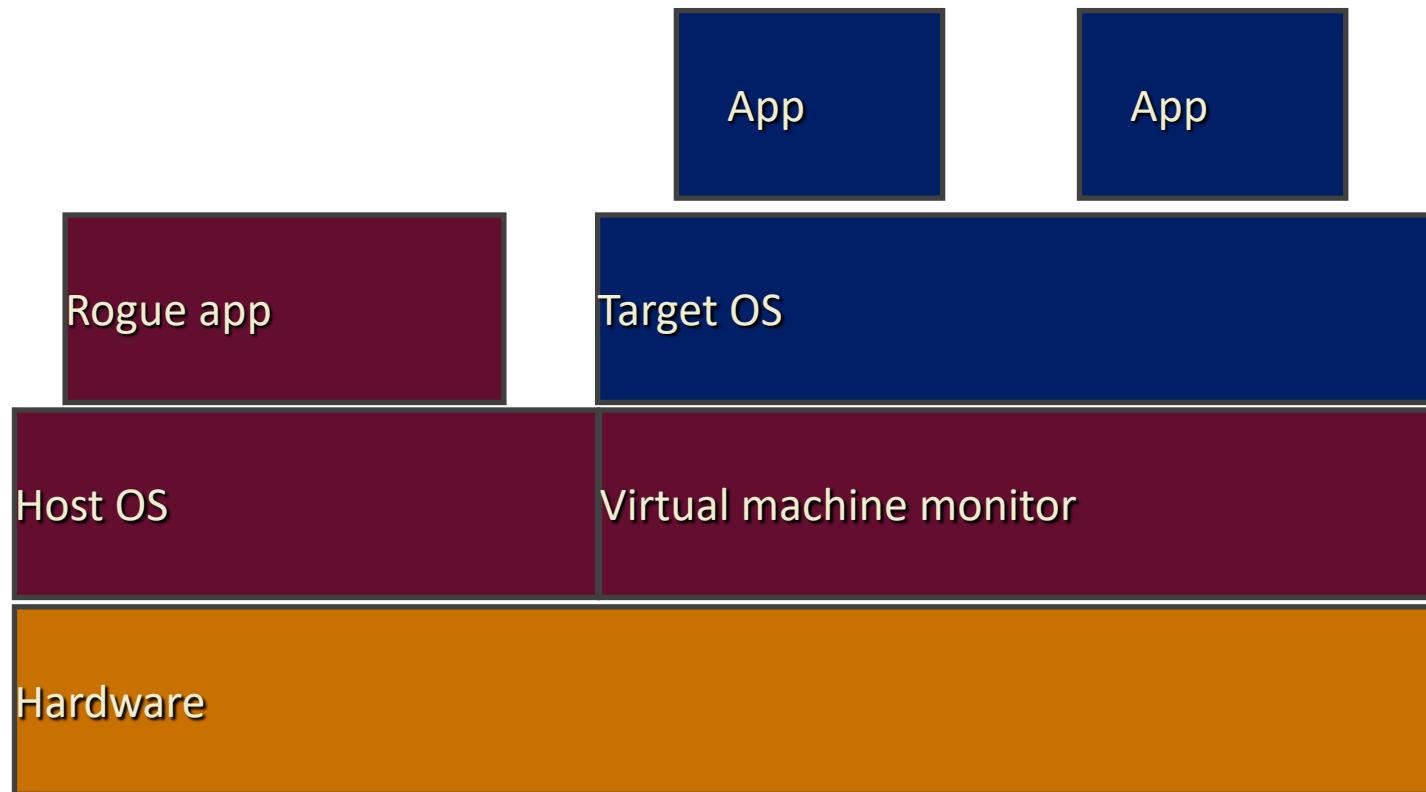
# Kernel rootkit



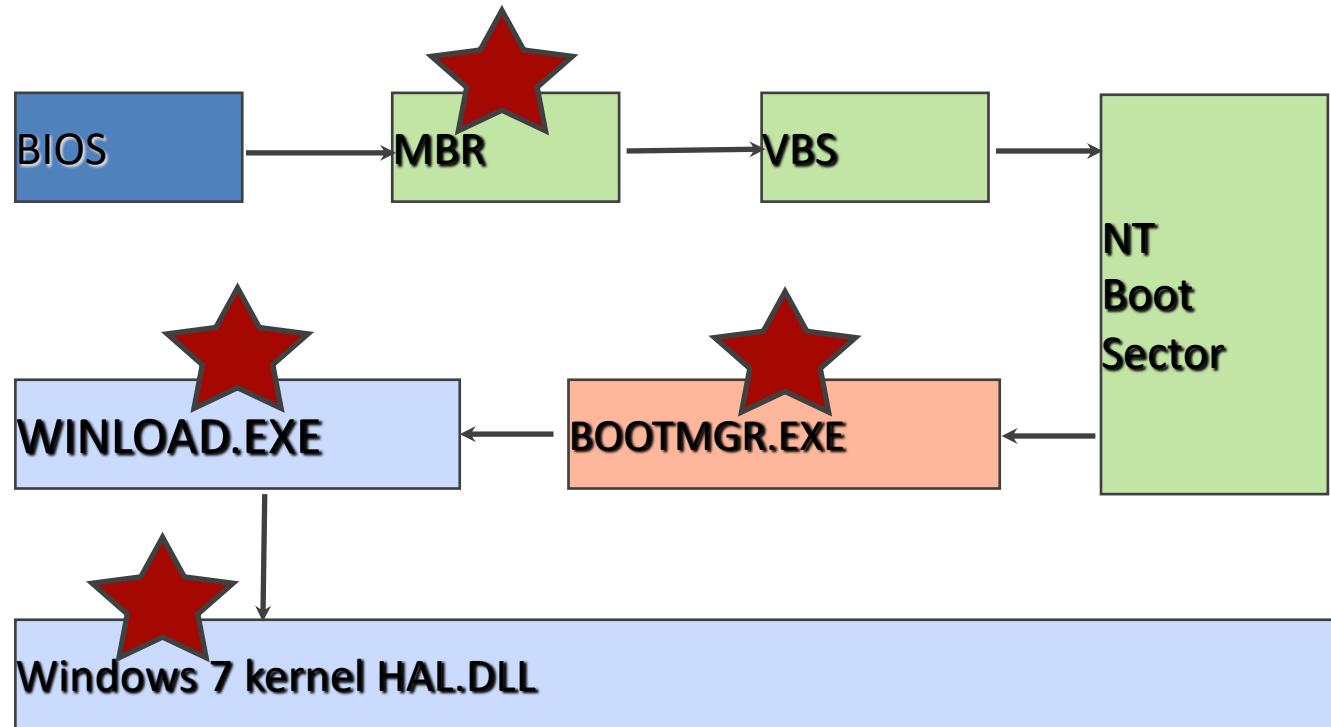
# Hypervisor rootkit



# Hypervisor rootkit



# MBR/Bootkit



All software is susceptible to attacks

How to run computation on data with  
confidentiality and integrity?

Make hardware the root of trust  
and  
imagine a world where ...

**Input File**

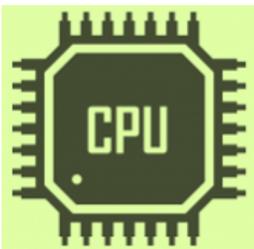
**Private Key**

## Signing Application



Bank's  
Web Server

## Operating System



```
4200 100e 028f 0e42 8e18 4503  
0e42 8c28 4405 300e 0686 0e41  
500e b803 0a01 380e 0e44 4130  
4220 180e 0e42 4210 080e 0b47  
004c 0000 c0ac 0000 a8c8 ffef  
4200 100e 028f 0e45 8e18 4203
```

```
"(private-key "  
"  (ecc "  
"    (curve Ed25519)"  
"    (flags eddsa)"  
"    (q #41ED58EA4FC9566FD510F357A426B4C25EB1859877D53CF4C19C43BD01F45A64#)"  
"    (d #B60569A2D9E566E39B99208A06BC9FB7B48F4BE1FD4BD0865C2423B6FCBBED39#)"  
"    )" "  
"  );";
```



Trusted



Untrusted

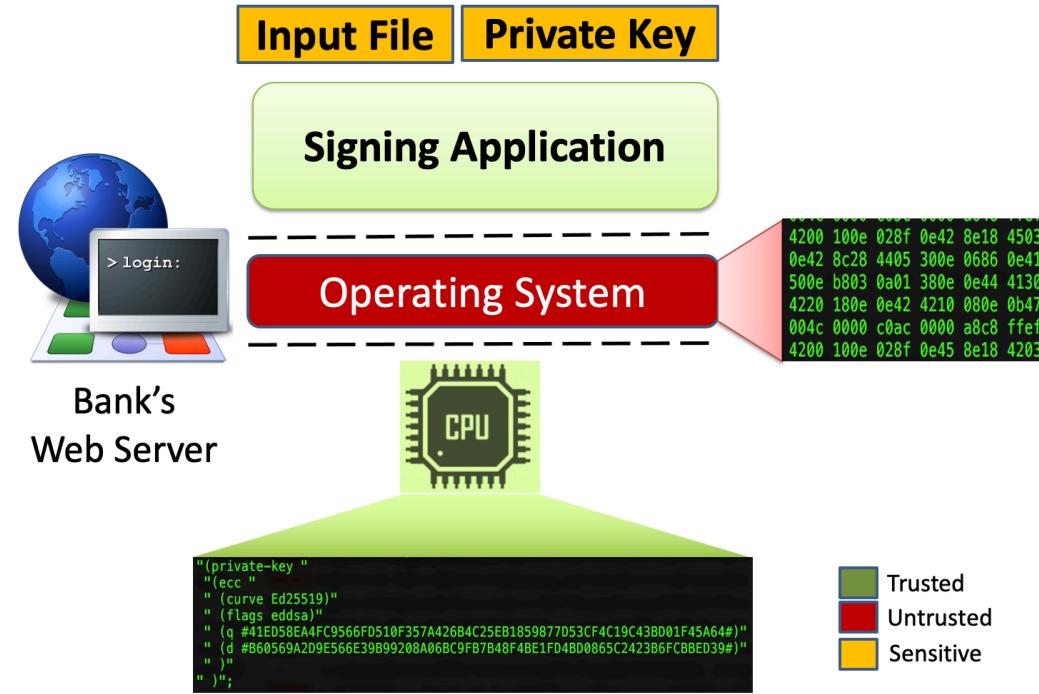


Sensitive

# Hurdles in achieving this dream

- How to bootstrap trust?
  - Ensure that the CPU is genuine (e.g., Intel Inside)
  - Ensure that we are running on a real machine (vs virtualized)
- How and where to store secrets if everything is malicious?
- What about attacks during execution?
- How to load an application if the OS is malicious?
- What about OS services?

# Break for questions

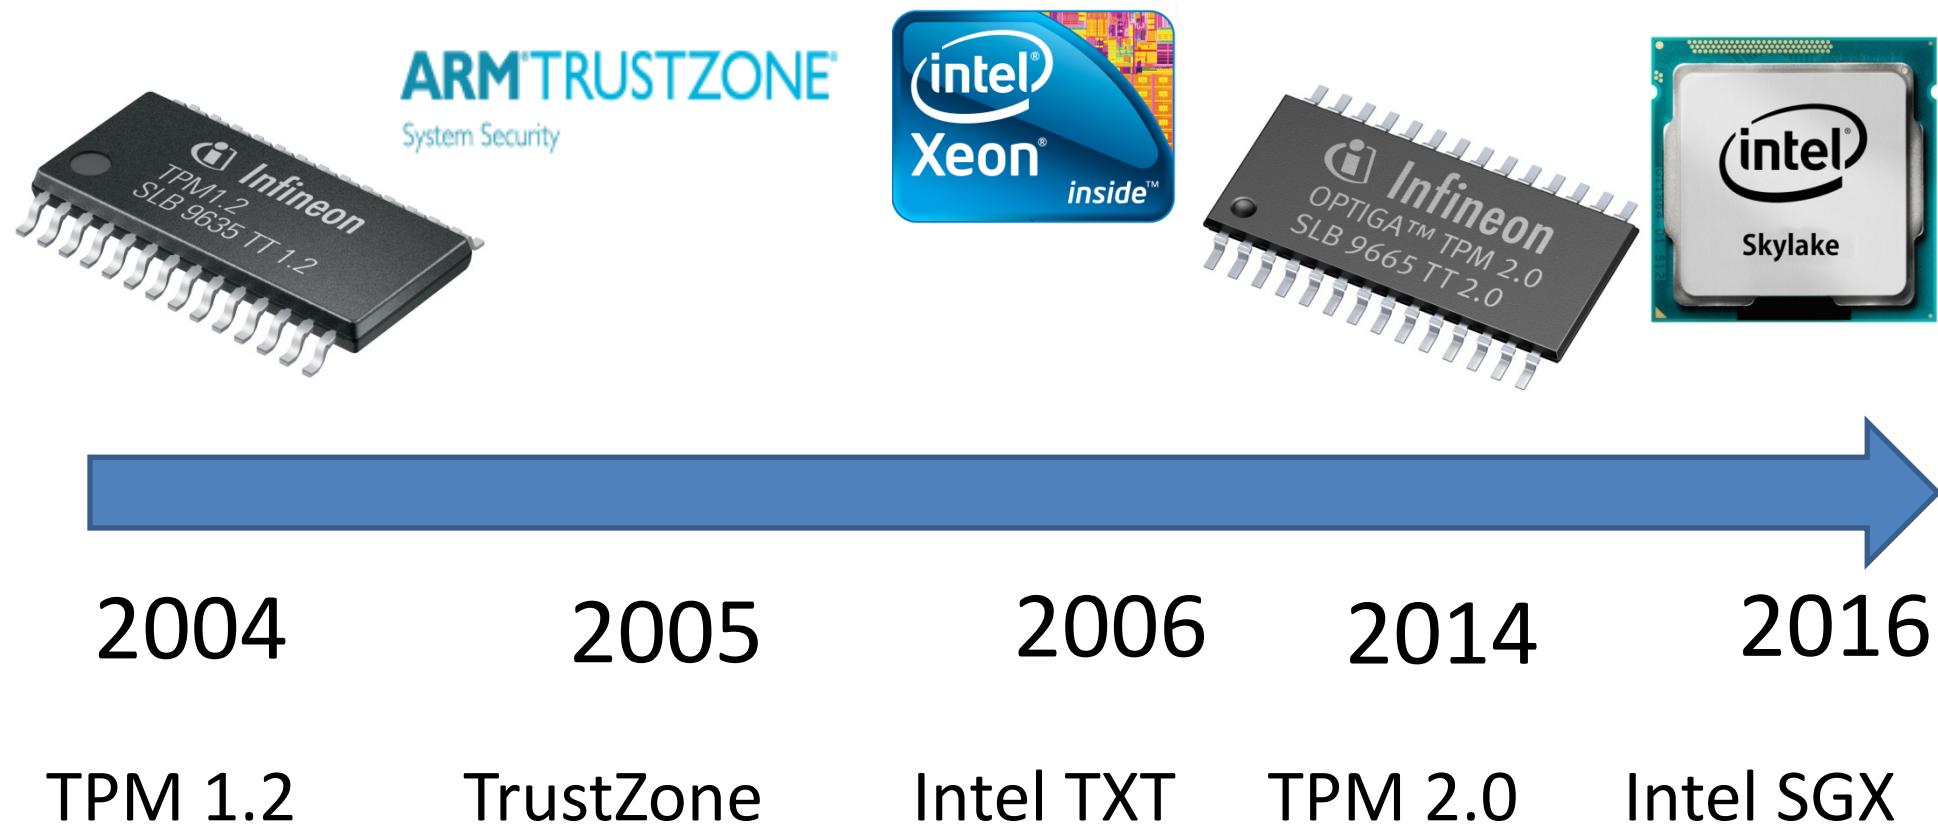


This is possible (yay!) if we design things carefully\*

# Trusted Computing Hardware

- A hardware root-of-trust vs. assumed software trust
- Why trust hardware?
  - Tamper-resistant from all software
  - Highest level of privilege, can monitor everything
  - Slow to iterate but can be verified

# Evolution of Trusted Computing



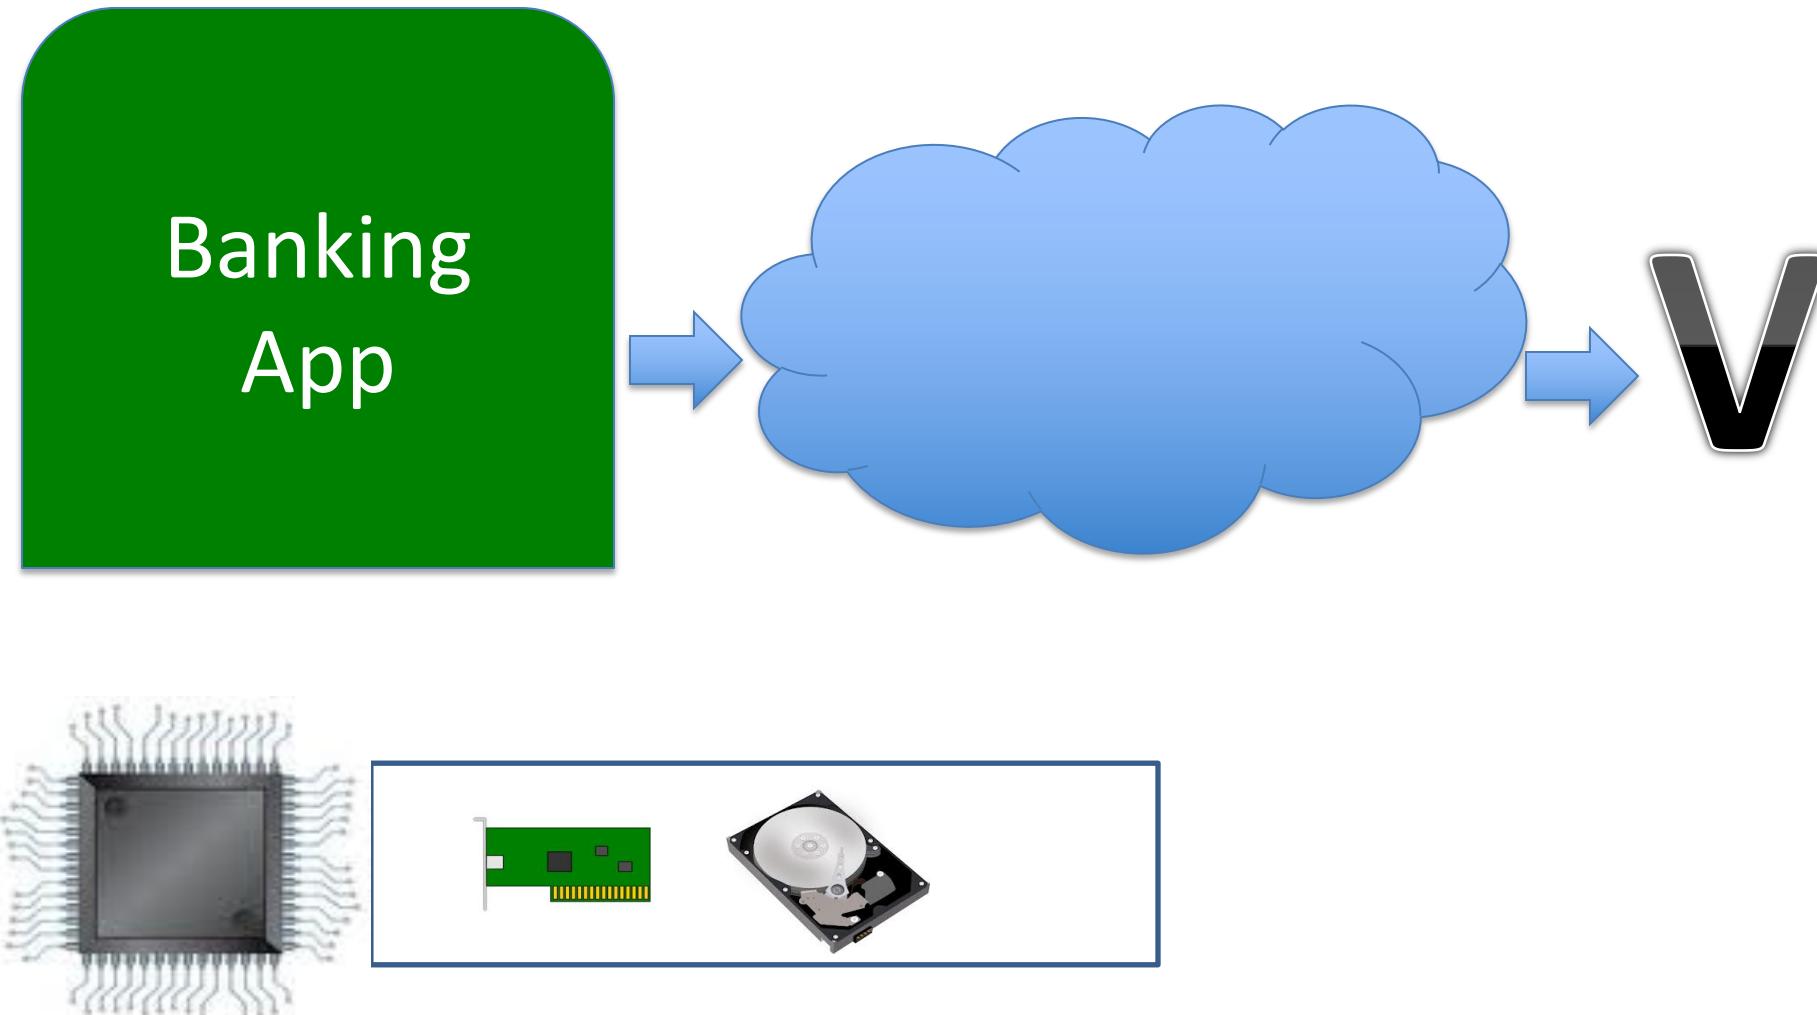
# Trusted Computing Primitives

- Remote Attestation
- Trusted Boot
- Sealed Storage
- Isolated Execution

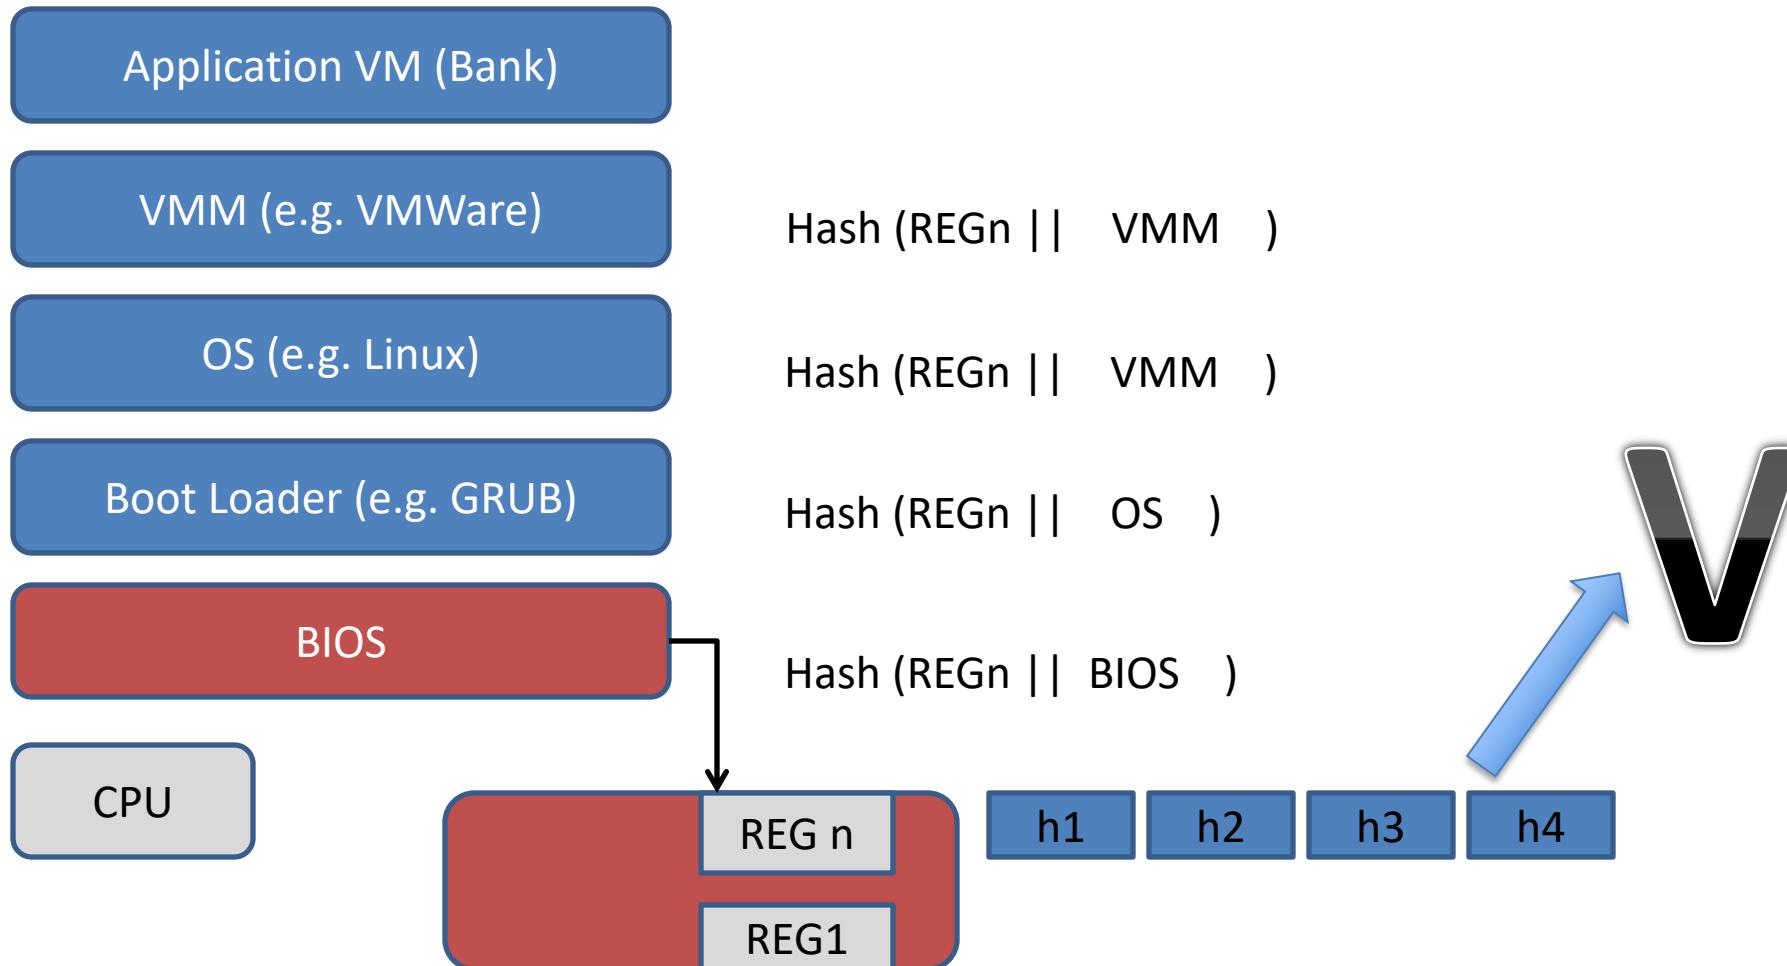
# Trusted Computing Primitives (I): Remote Attestation

- **Goal**: Prove to remote party what software is loaded on my machine

# Trusted Computing Primitives (I): Remote Attestation



# Trusted Computing Primitives (I): Remote Attestation

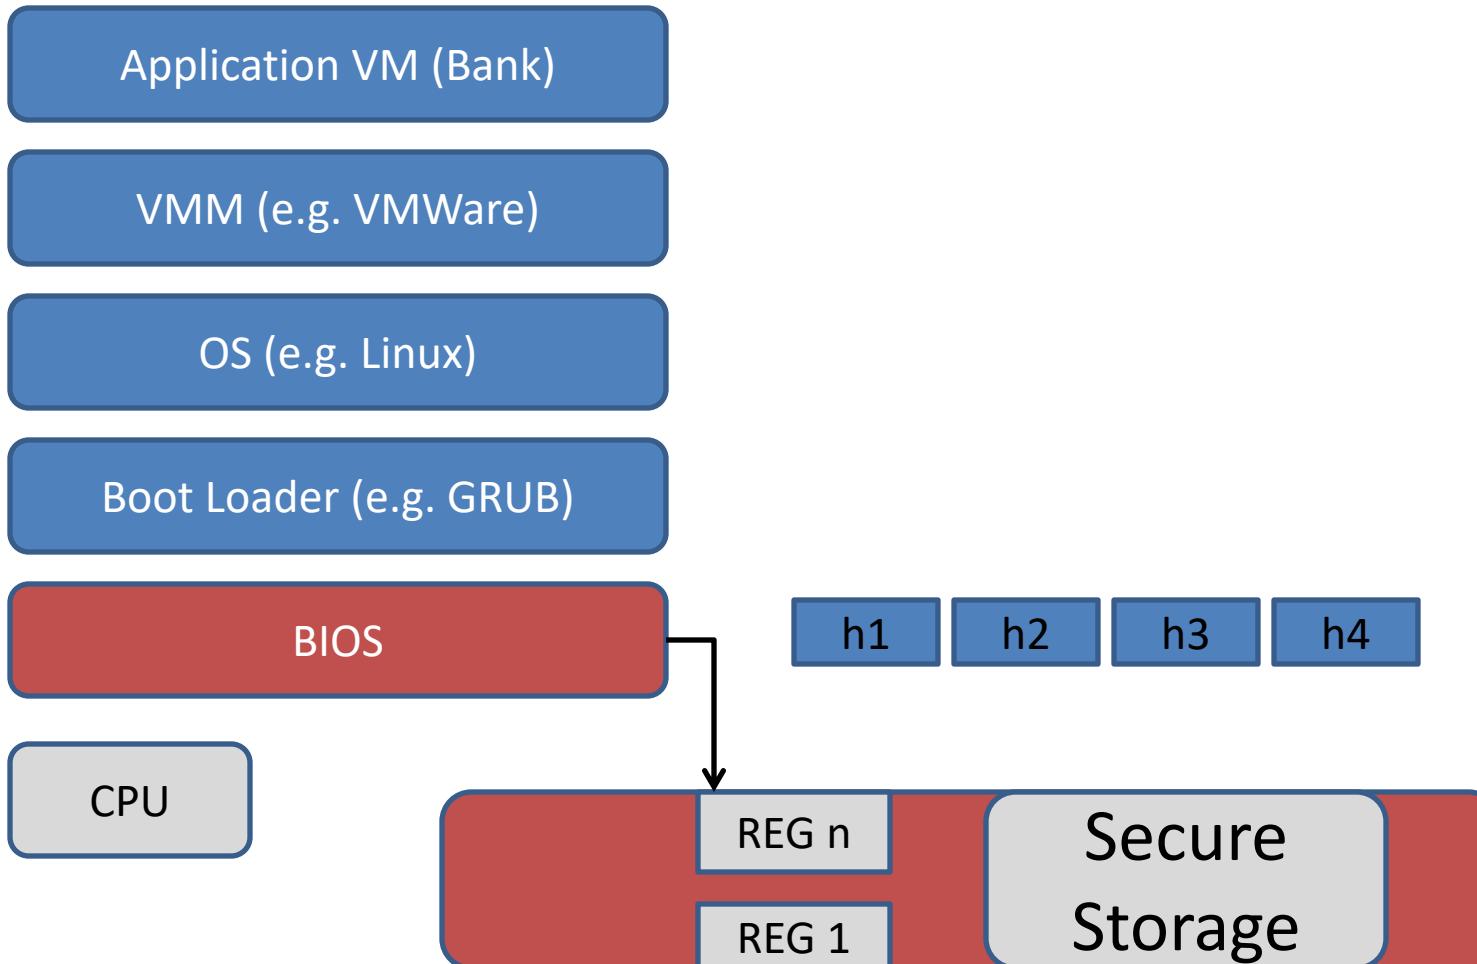


Each CPU has a signing key

# Trusted Computing Primitives (II): Trusted Boot

- **Goal:** To perform a measured launch of an OS kernel/VMM
- **Example:** Bank allows money transfer only if customer's machine runs an “up-to-date” OS with patches

# Trusted Computing Primitives (II): Trusted Boot

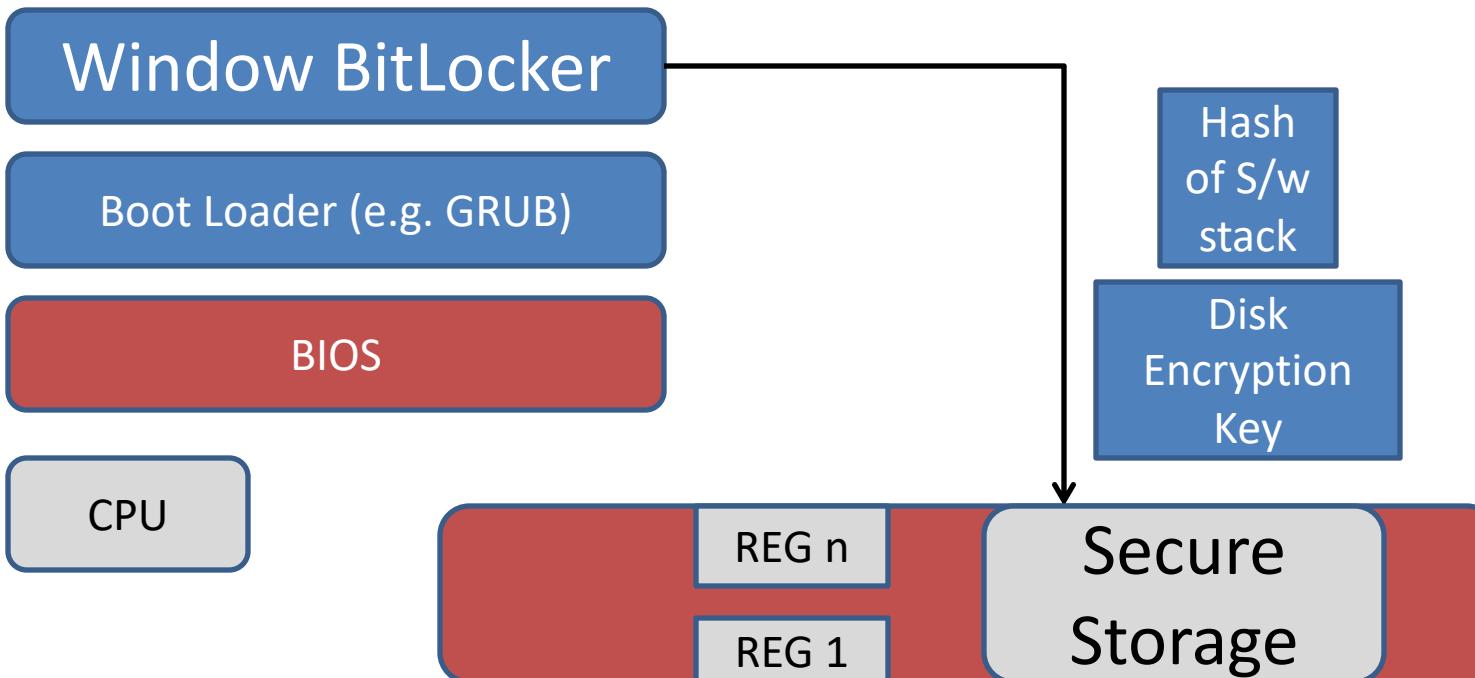


# Trusted Computing Primitives(III): Sealed Storage

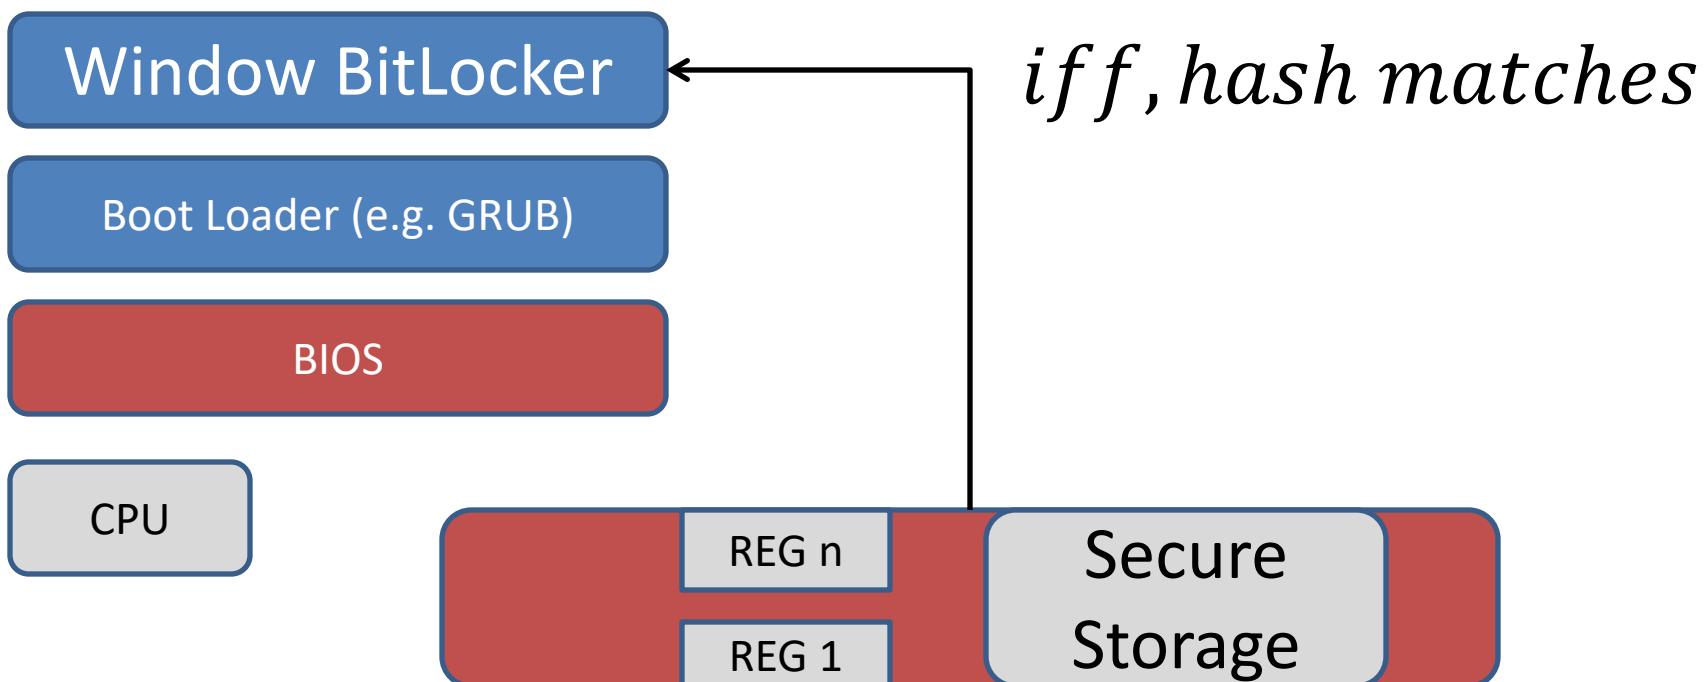
- **Goal:** Protect private information by binding it to platform.  
Data can be released only to a particular combination of SW & HW

# Trusted Computing Primitives(III): Sealed Storage

Use CPU Measurements & Secure Storage for  
Disk Encryption Systems?



# Trusted Computing Primitives (III): Sealed Storage



# Use Case: Full Volume Encryption

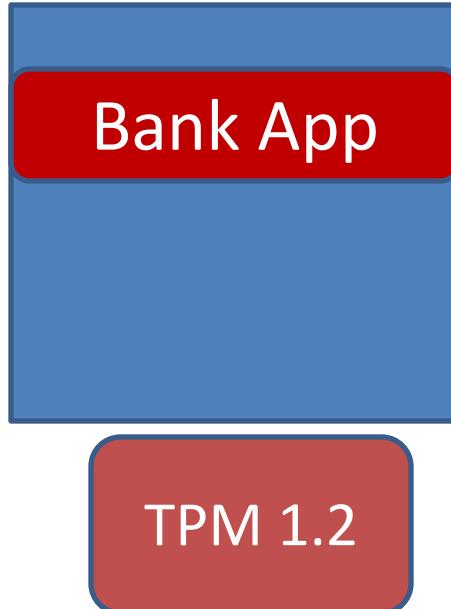
- Encryption at the block level underneath file system
- Everything in the volume is encrypted.
- BitLocker was first used by Microsoft since Windows Vista
- BitLocker takes advantage of TEEs
  - Top level root key sealed in hardware
  - Root key encrypts disk encryption key, which encrypts sector data
- CPU protects disk encryption key by encrypting it
- CPU releases key only after comparing hash of early (unencrypted) boot files with previous hash

# Trusted Computing Primitives (IV): Isolated Execution

- Goal: Execute sensitive code in isolation from other malicious programs running on the same machine

# Isolated Execution with Late Launch / DRTM

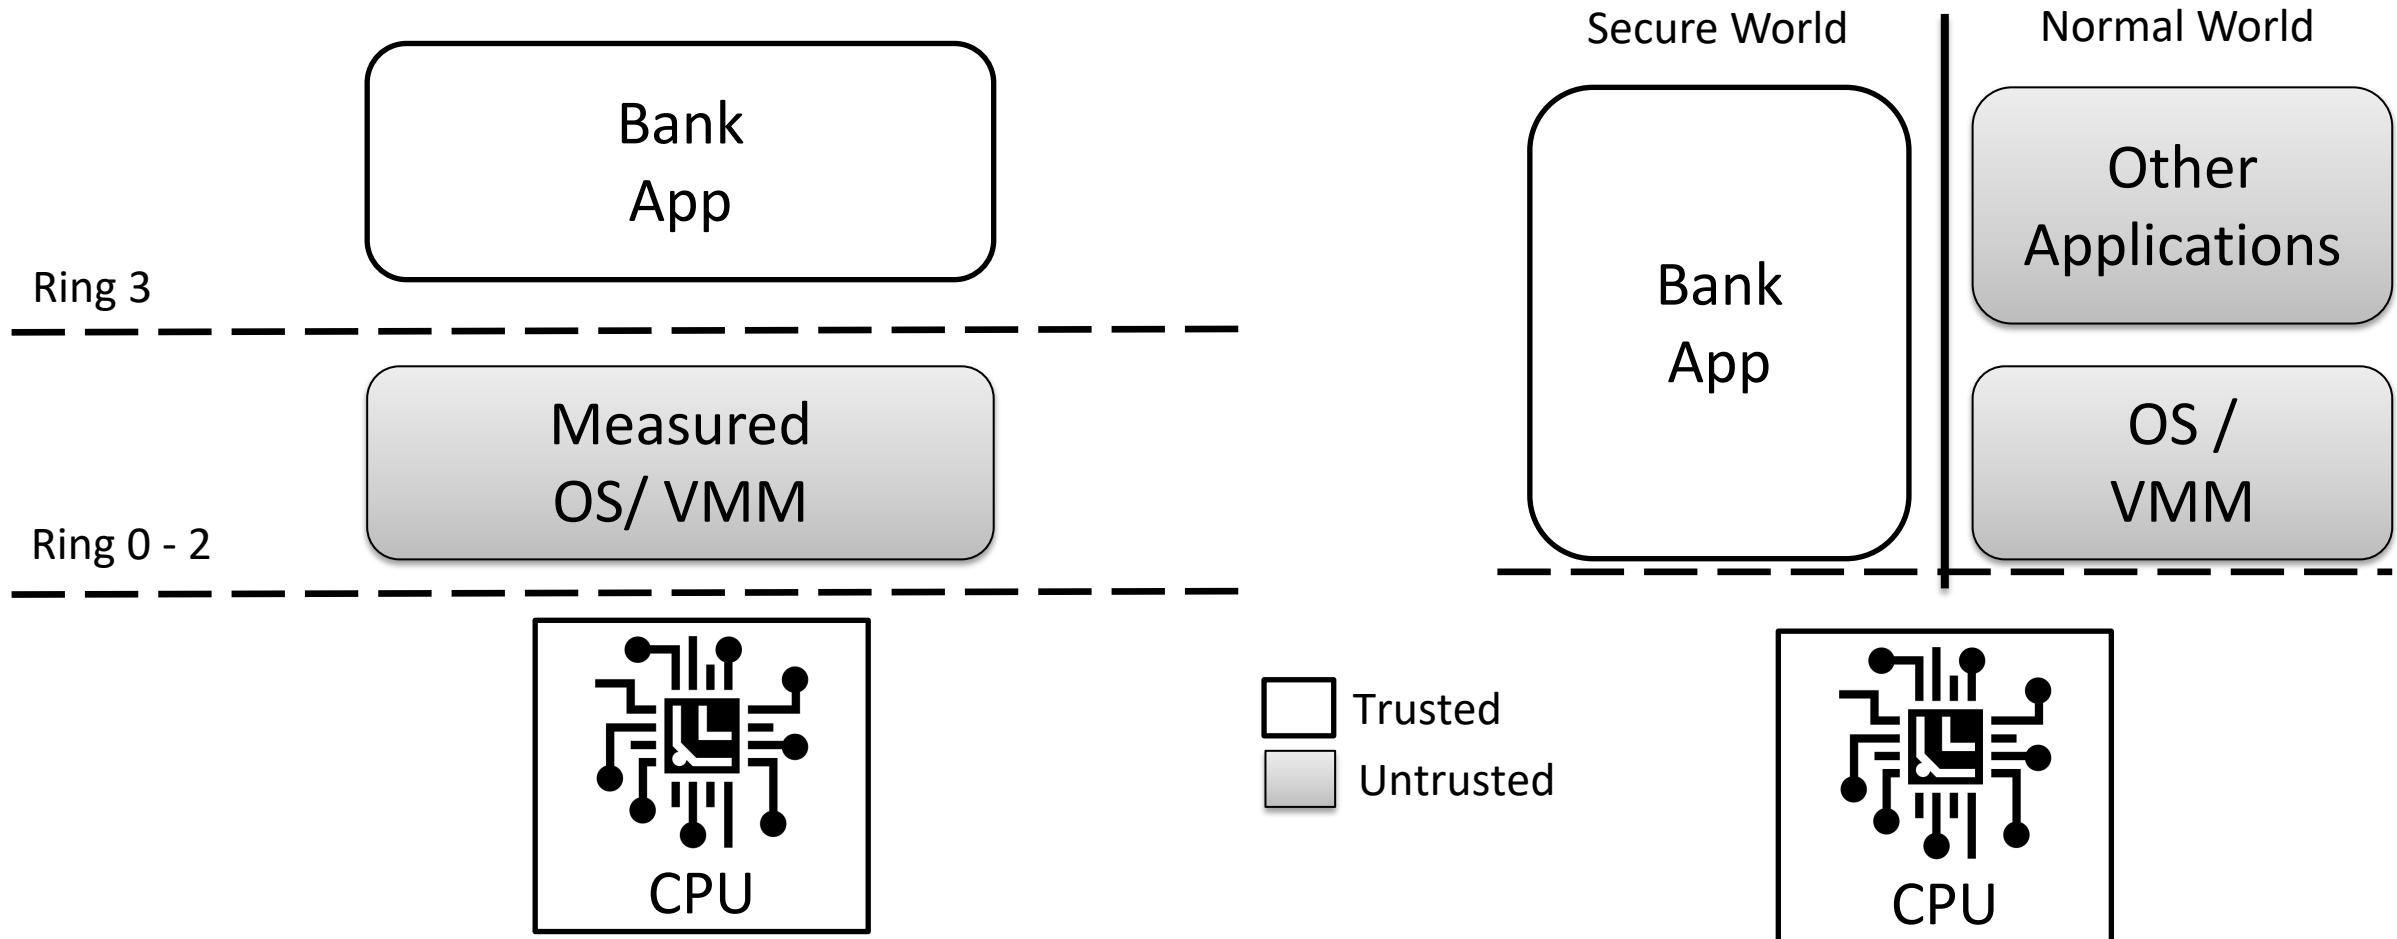
- One Approach: Intel TXT “SKINIT/SENTER”



On SENTER:

- TPM hashes memory region in PCR17
- Checks if valid
- Executes the trusted region code

# Isolated Execution Models

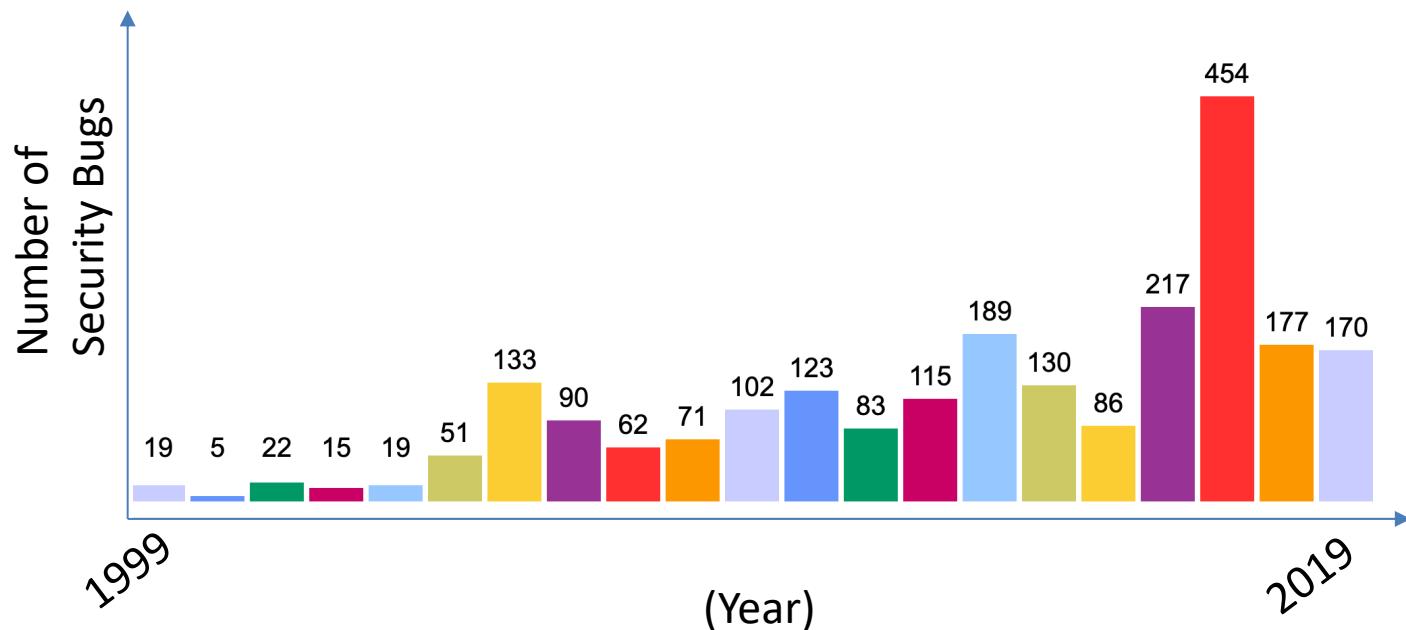


# Break for questions

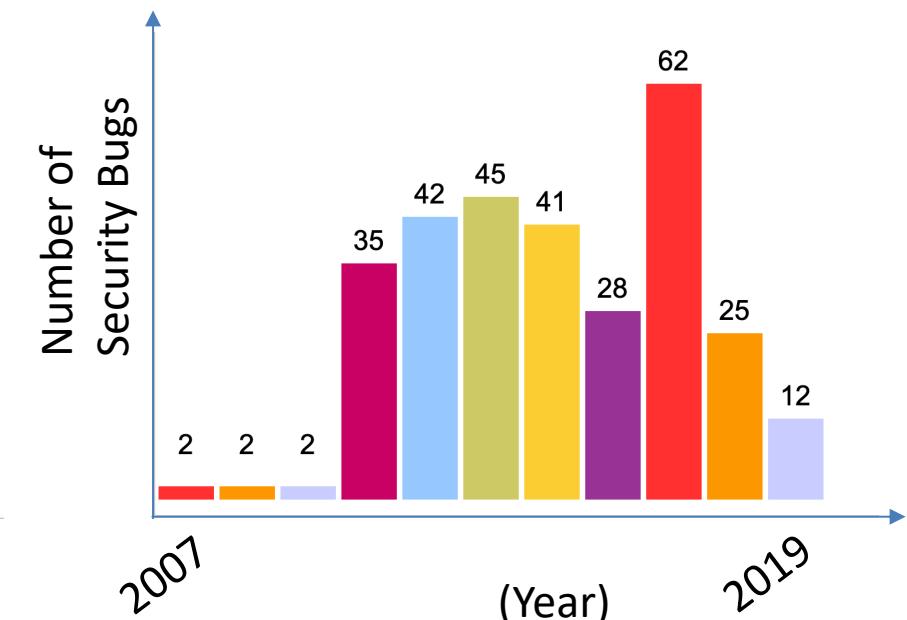
# How much code are we trusting?

We measure the OS but still assume it to be bug-free

Operating System (Linux Kernel)  
27 Million Lines

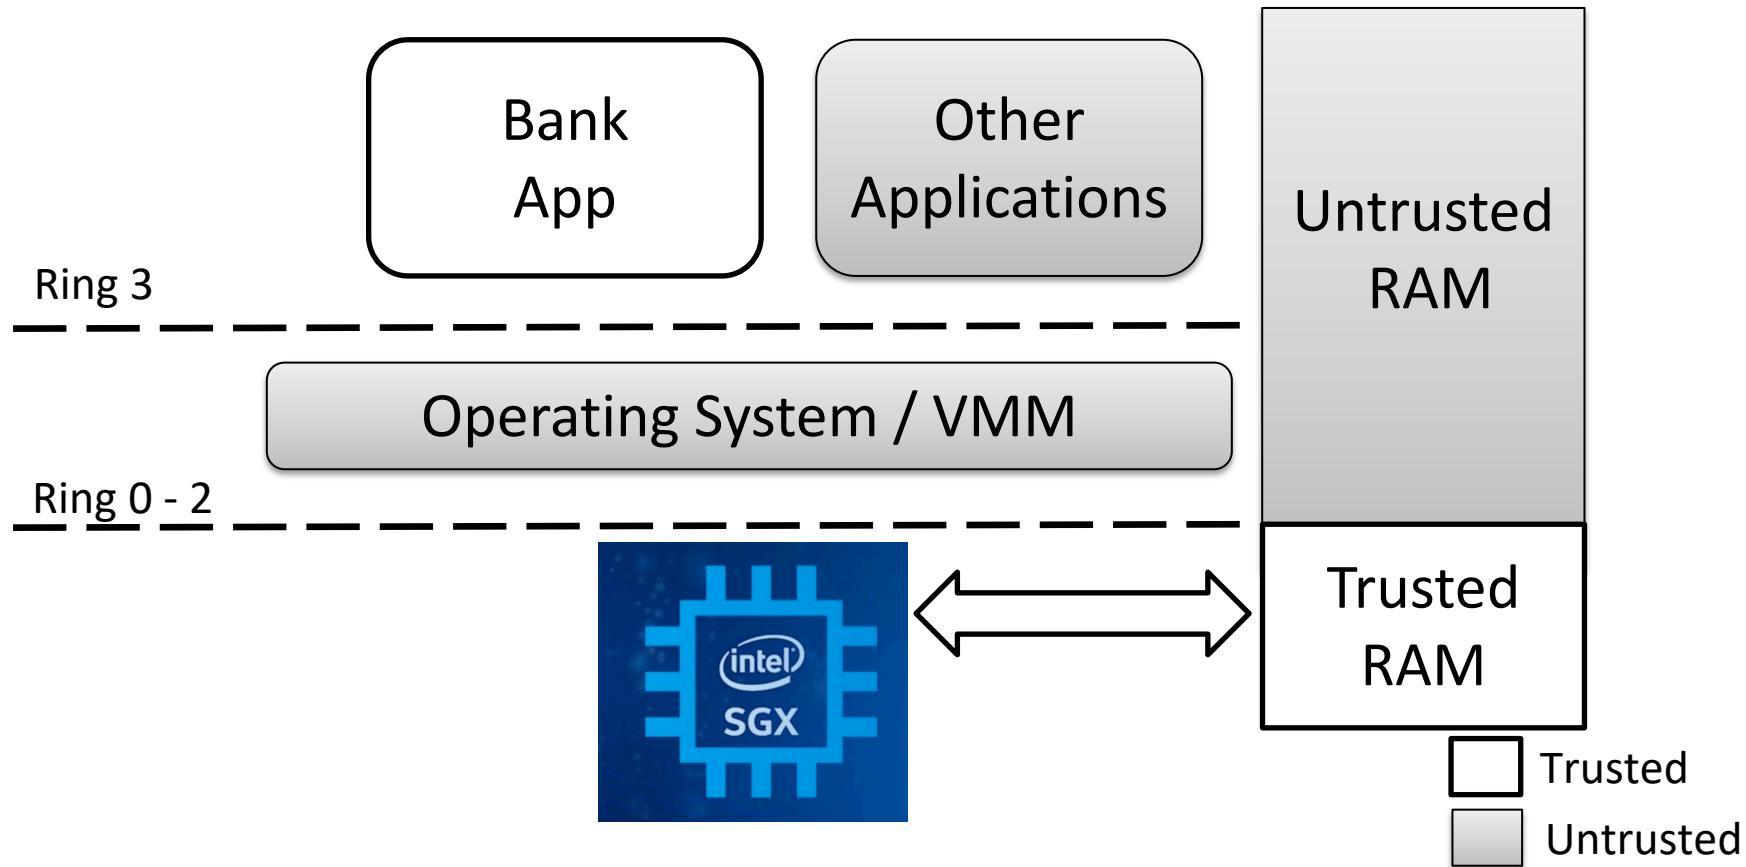


Hypervisor (XEN)  
0.5 Million Lines

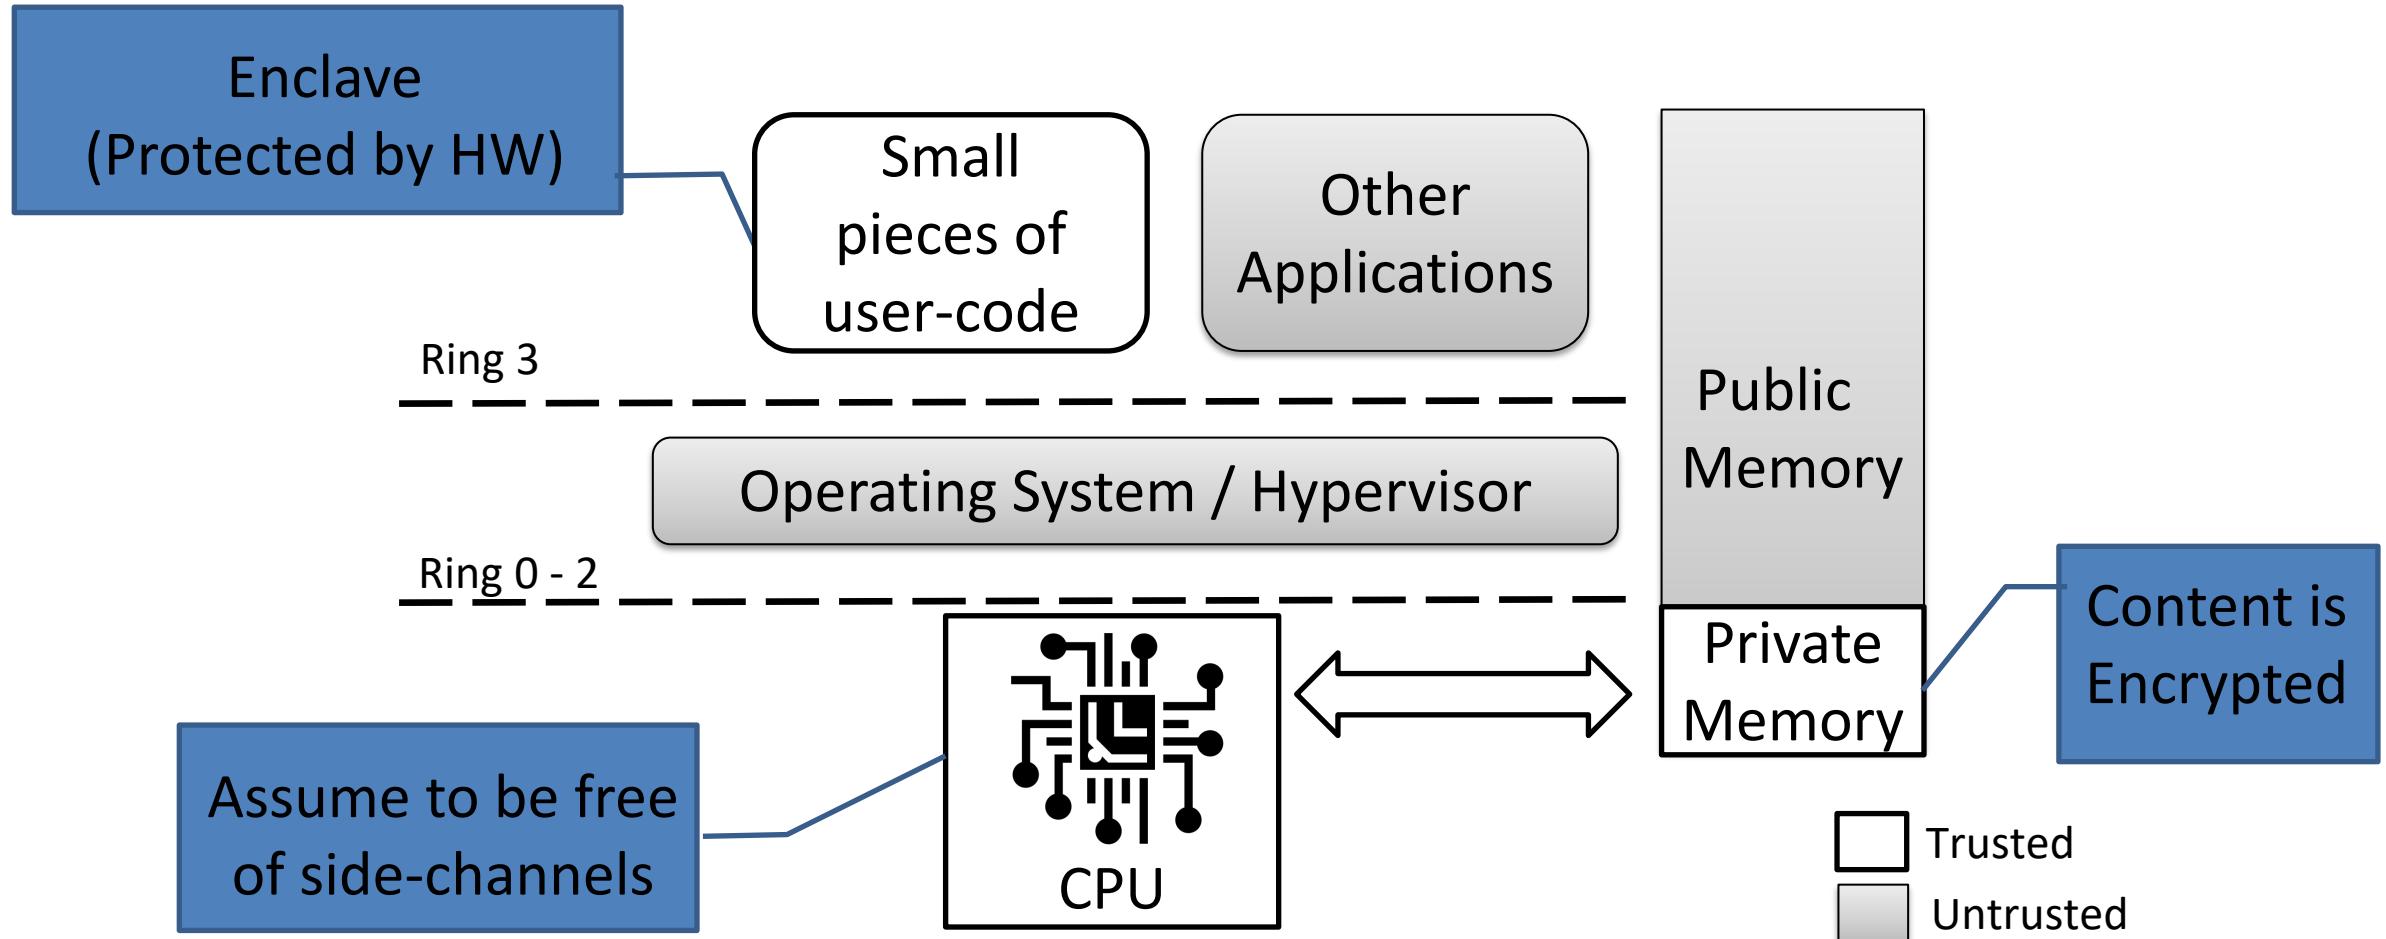


Another approach for achieving  
Isolated Execution

# SGX Isolated Execution



# Intel SGX Enclaves



# Design Choice: Untrusted OS/VMM

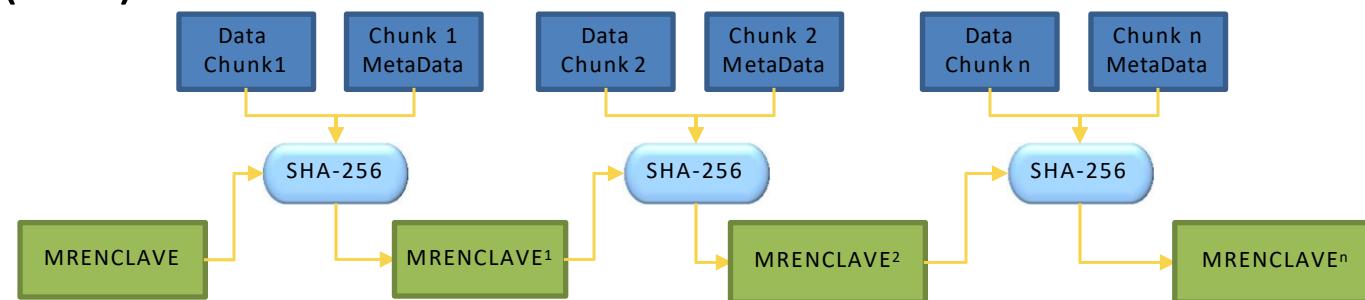
# Enclave Measurement

When building an enclave, Intel® SGX generates a cryptographic log of all the build activities

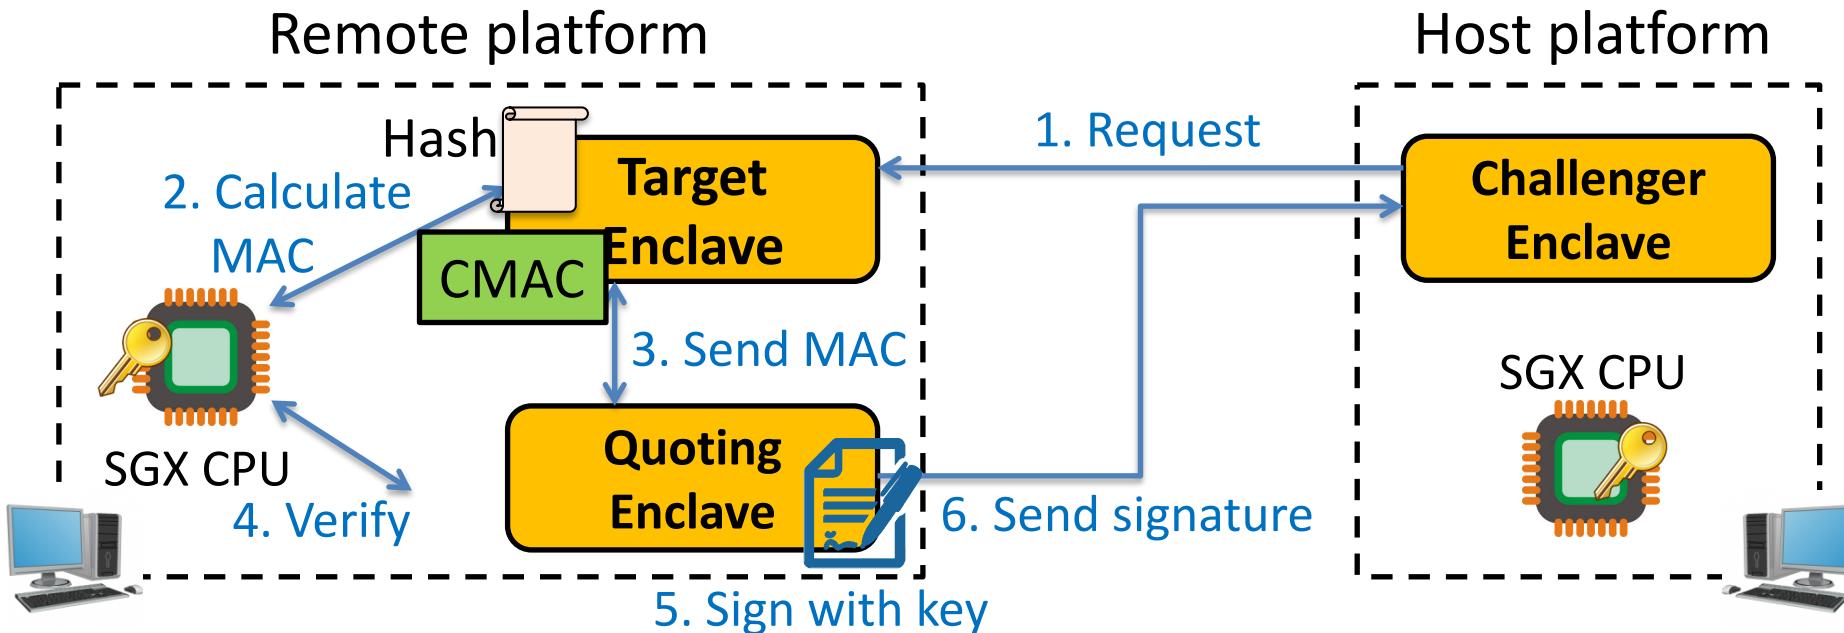
- Content: Code, Data, Stack, Heap
- Location of each page within the enclave
- Security flags being used
- Order in which it was built

**MRENCLAVE** (“Enclave Identity”) is a 256-bit digest of the log

- Represents the enclave’s software trusted computing base (TCB)



# SGX Remote Attestation

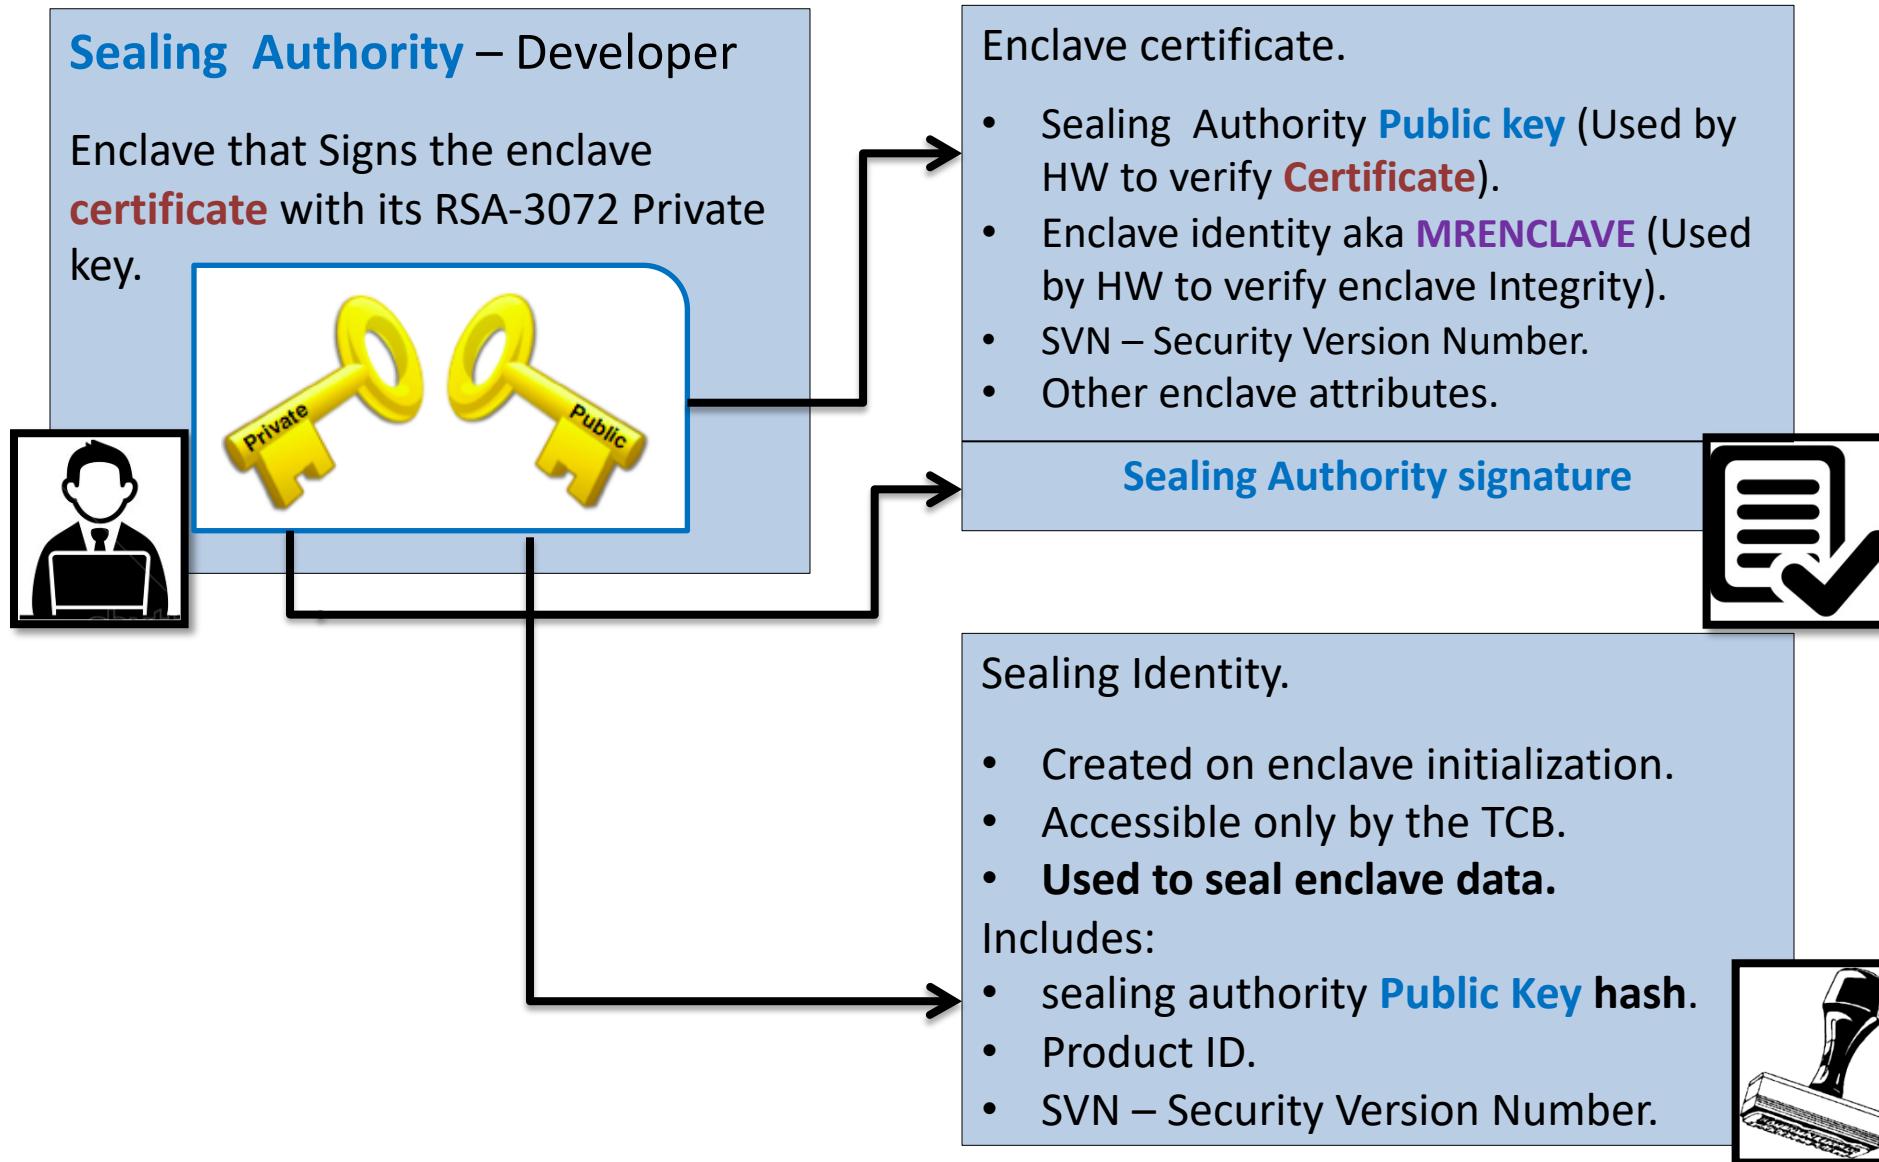


- Attest an application on remote platform
- Check the identity of enclave (**hash of code/data pages**)
- Can establish a “**secure channel**” between enclaves

# SGX Sealed Storage

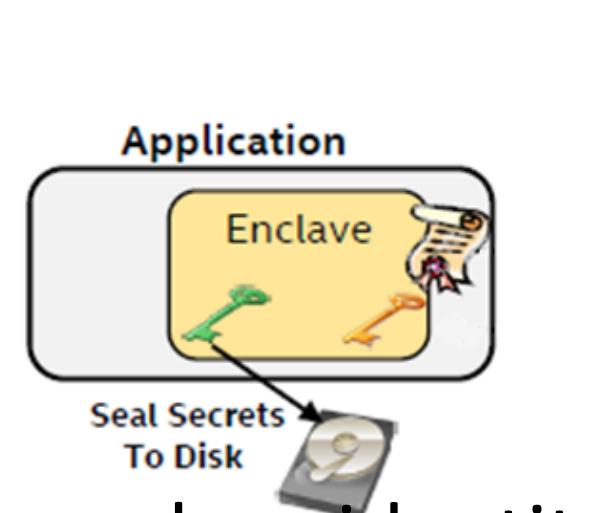
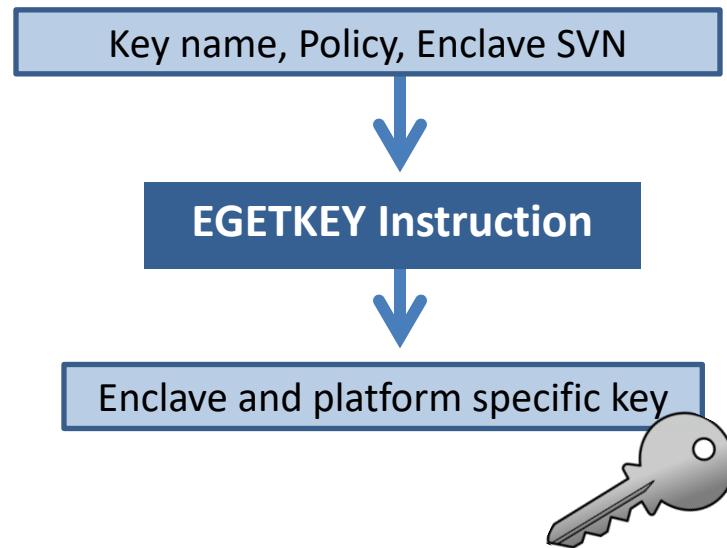
- Cryptographically protect data when it is stored outside enclave

# SGX Sealing Actors



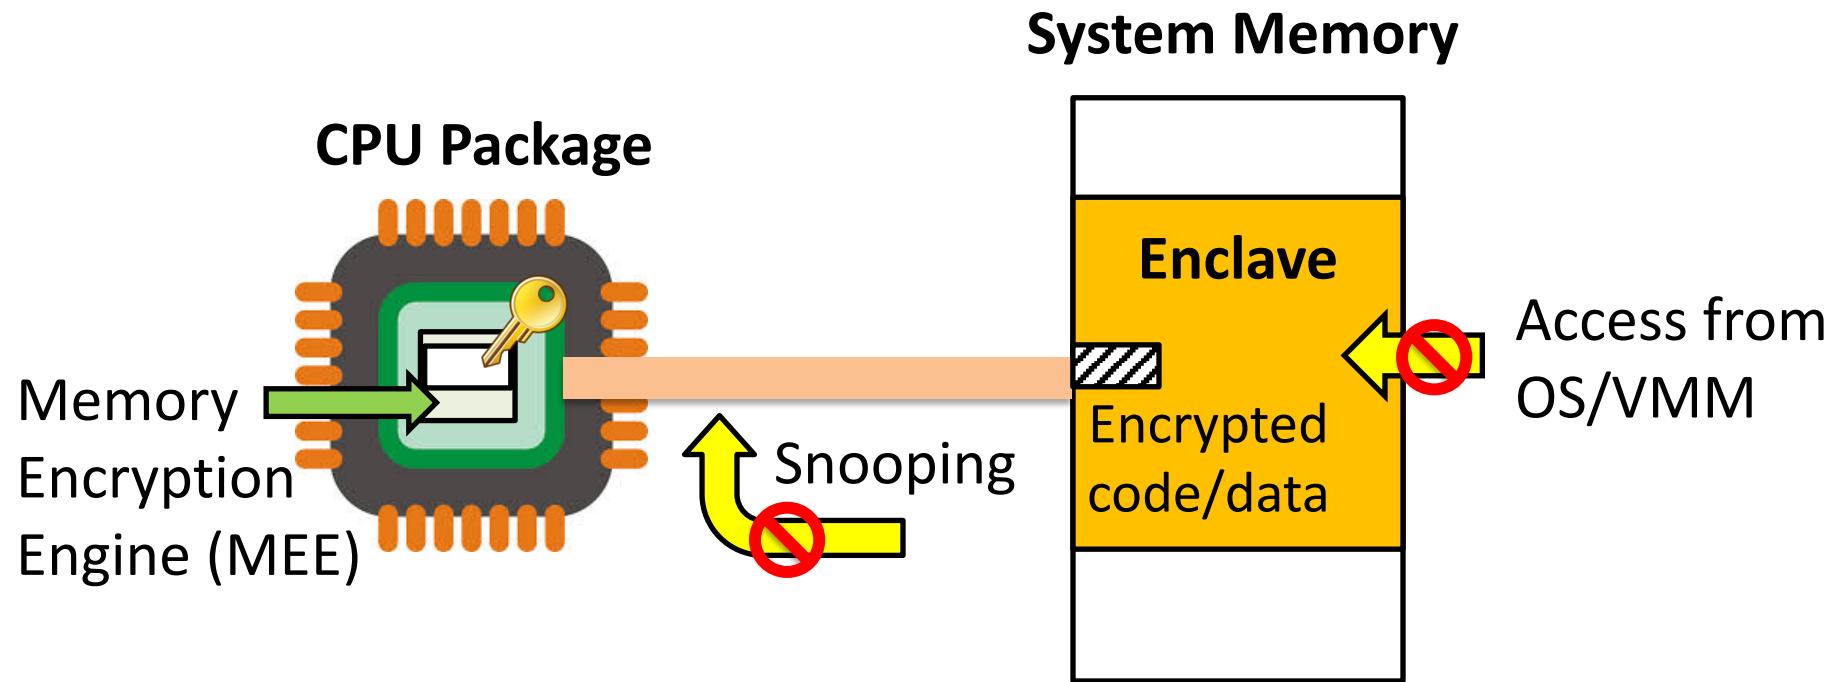
# SGX Sealing Process

- 1. Enclave calls EGETKEY
- 2. Sets sealing key
- 3. Seals data with key
- 4. Writes sealed data to untrusted storage



Data sealed under processor key, enclave identity...

# SGX: Isolated Execution



- Application keeps its data/code inside the “enclave”
  - Smallest attack surface by reducing TCB (App + processor)
  - Protect app’s secret from untrusted privilege software (e.g., OS, VMM)

# Architecture Overview

## Enclave

Trusted Execution Environment embedded in application

Provides confidentiality and/or integrity of its own code/data

## EPC (Enclave Page Cache)

Trusted Memory for enclave

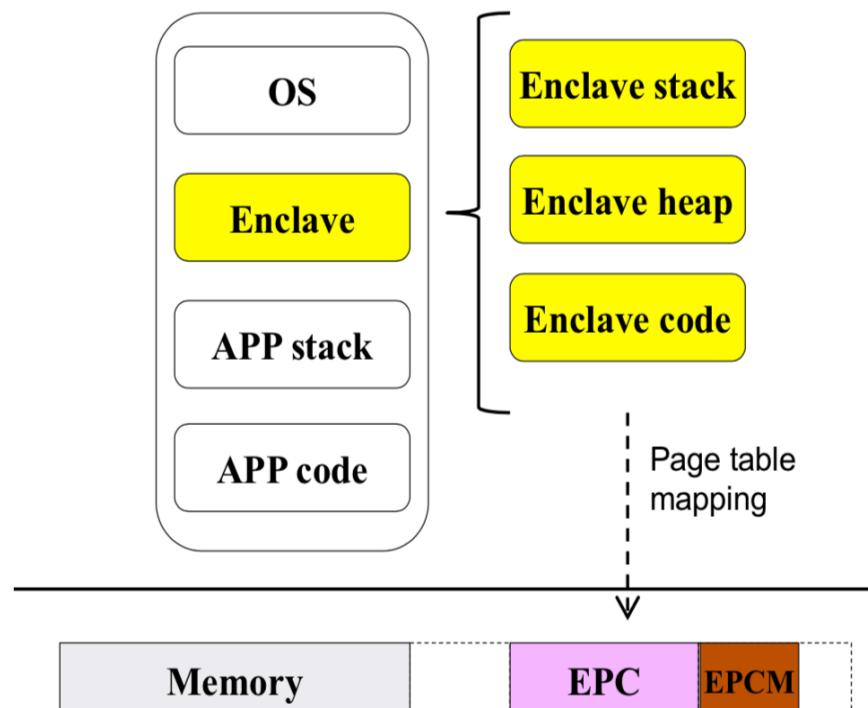
Typically reserved by BIOS

Limited size (ex, 32/64/128M)

## EPCM (Enclave Page Cache Map)

Used by HW to track EPC

(not-visible to SW)

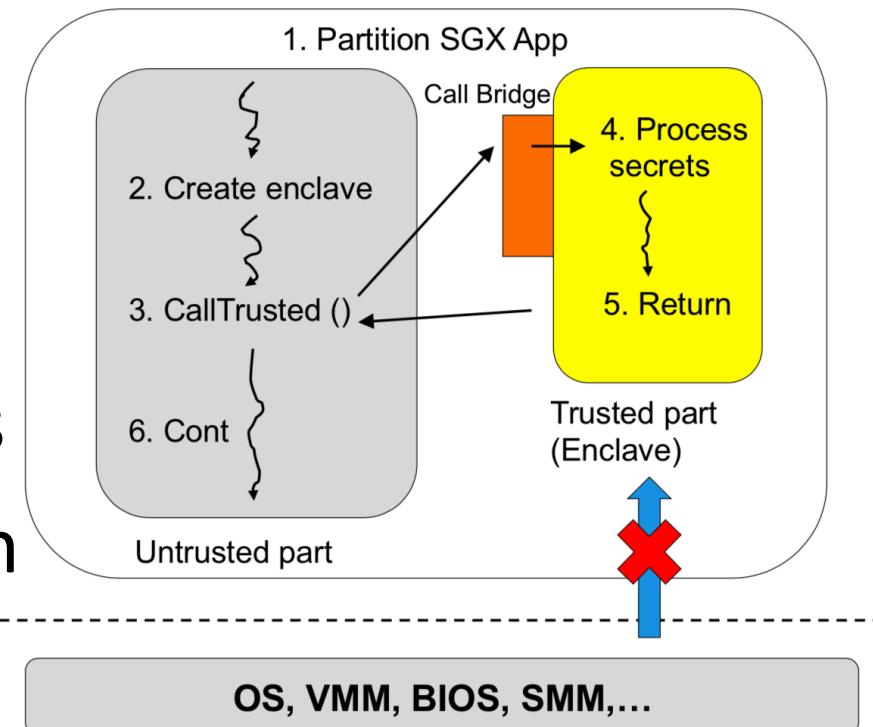


# Design Choice: Untrusted OS

- Against the typical setting used for decades
- What should we protect?
  - **Creation of the enclave**
  - **Code & data inside the enclave**
  - **Data generated by the enclave**
  - **Services like timers, random numbers**
  - Start and end of enclave
  - Enclave-enclave communication
  - Memory management of the enclave
  - IO & System calls by the enclave

# Controlled Entry & Exit Points

- Define & partition app
  - untrusted and trusted
- App creates enclave
- Calls a trusted function
- Enclave code executes
- Trusted function returns
- App continues execution



**Input File**

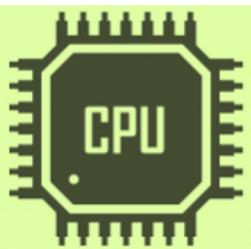
**Private Key**

## Signing Application



Bank  
Web Server

## Operating System



```
4200 100e 028f 0e42 8e18 4503  
0e42 8c28 4405 300e 0686 0e41  
500e b803 0a01 380e 0e44 4130  
4220 180e 0e42 4210 080e 0b47  
004c 0000 c0ac 0000 a8c8 ffef  
4200 100e 028f 0e45 8e18 4203
```

```
"(private-key "  
"  (ecc "  
"    (curve Ed25519)"  
"    (flags eddsa)"  
"    (q #41ED58EA4FC9566FD510F357A426B4C25EB1859877D53CF4C19C43BD01F45A64#)"  
"    (d #B60569A2D9E566E39B99208A06BC9FB7B48F4BE1FD4BD0865C2423B6FCBBED39#)"  
"    )" "  
"  );";
```



Trusted



Untrusted



Sensitive

# Applications of TEEs

- Server-side apps
  - webserver, database
- Client-side apps
  - hardware security module (HSM)
  - Enterprise Rights Management (DRM)
  - Faster encrypted computation
- Securing remote execution
  - Data analytics (e.g., ML) in the cloud
  - Distributed computing (e.g., Tor, consensus protocols)

# BREAK

Are TEEs (e.g., SGX) sufficient?

# Covert Channels

- Definition: “An unintended channel of communication between two untrusted programs”
- E.g., Shared Cache Latency
  - Sender
    - Send bitval 1: Perform random memory access
    - Send bitval 0: Do nothing
  - Receiver
    - Rcv bitval 1: If long read time for a fixed memory location
    - Rcv bitval 0: If short read time for a fixed memory location
- Many traditional channels: Caches (all levels), Timing, Power, Disk, I/O, virtualization latency

# Side Channels

- Definition: “An channel of information leakage from an trusted to an untrusted program”
- E.g. Learning encryption keys of co-located VMs
- $E := x^e \pmod{N}$

```
SquareMult( $x, e, N$ ):  
    let  $e_n, \dots, e_1$  be the bits of  $e$   
     $y \leftarrow 1$   
    for  $i = n$  down to 1 {  
         $y \leftarrow \text{Square}(y)$   
         $y \leftarrow \text{ModReduce}(y, N)$   
        if  $e_i = 1$  then {  
             $y \leftarrow \text{Mult}(y, x)$   
             $y \leftarrow \text{ModReduce}(y, N)$   
        }  
    }  
    return  $y$ 
```

(S)  
(R)  
(M)  
(R)

Timing Differences

Cache Access Patterns

# Typical Defenses

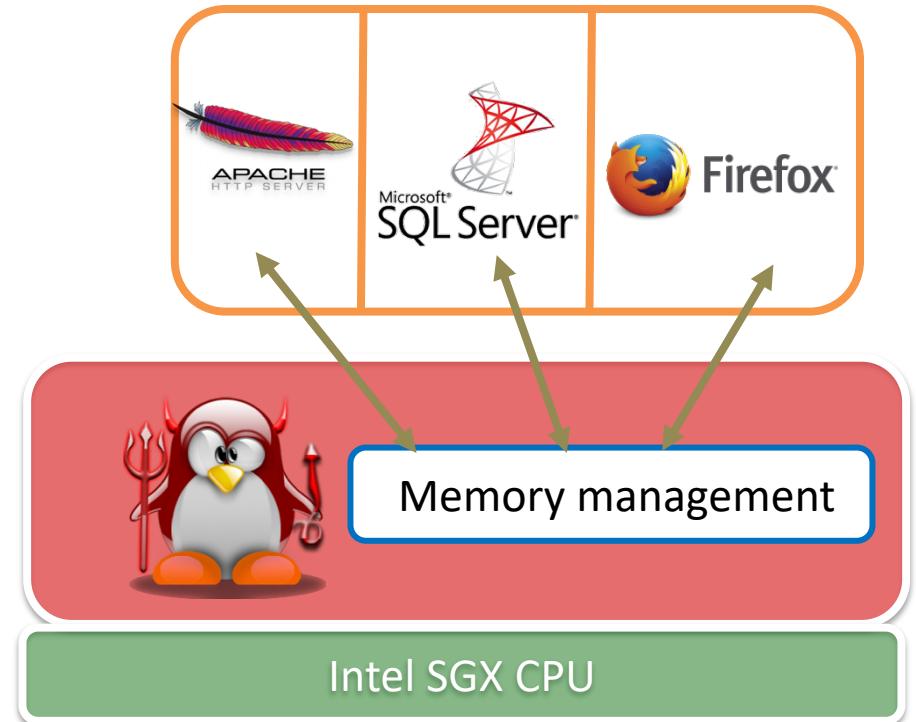
- Decouple the relationship between secrets and observations
- Determinism: Pad to normalize the channel for
- Indistinguishability: Add jitter, randomization, or noise
- Clear the state every time irrespective of the secret
- Disallow access to the channel

# Threat to TEE Guarantees

- TEEs do not directly protect enclaves against side channels ☹
- Enclaves may leak information via traditional channels
- Added Threat:
  - Recent channels via shared microarchitecture (e.g., spectre):
  - private caches, branch predictors, TLBs
- Leaking the CPU key has larger ramifications
  - e.g., Compromise remote attestation
- More bad news: TEEs create new side-channels

# Untrusted Memory Manager

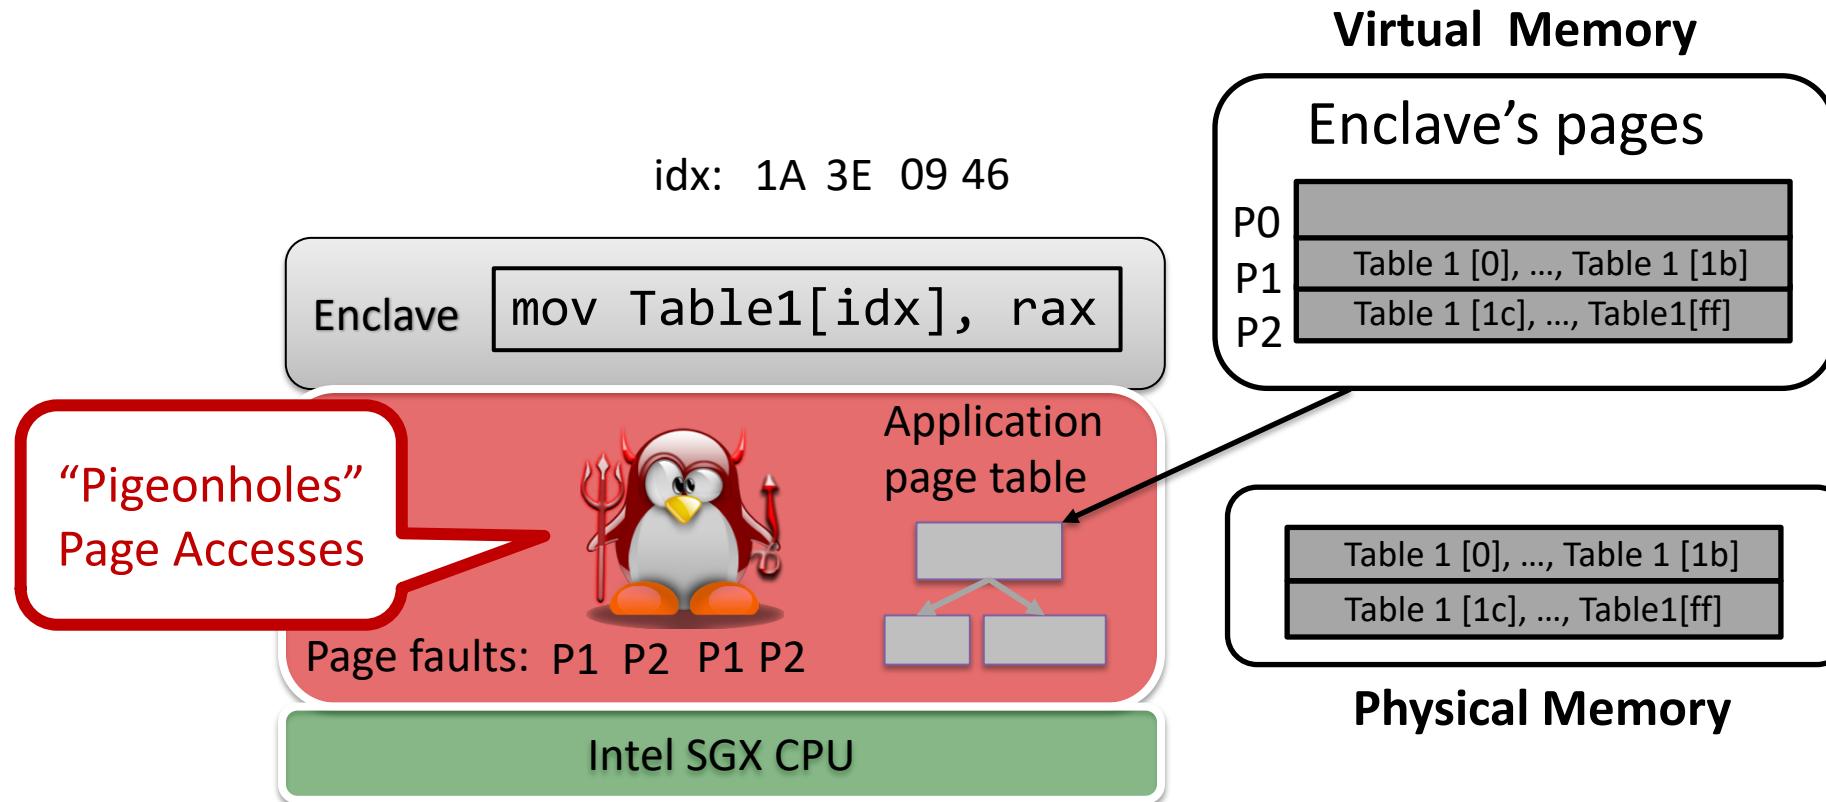
- OS manages all enclave resources
- New sources of side channels:  
system calls, file I/O
- Recent attack channels: Page faults



# Page-fault side-channel in TEEs

- Observation: The untrusted OS controls the memory management of the enclave
- OS decides how much memory to allocate for each enclave
- OS is notified of memory events (e.g., page faults)
  
- In SGX, the page fault reports only the page number, not the offset within the page

# Attack Example I: AES Encryption

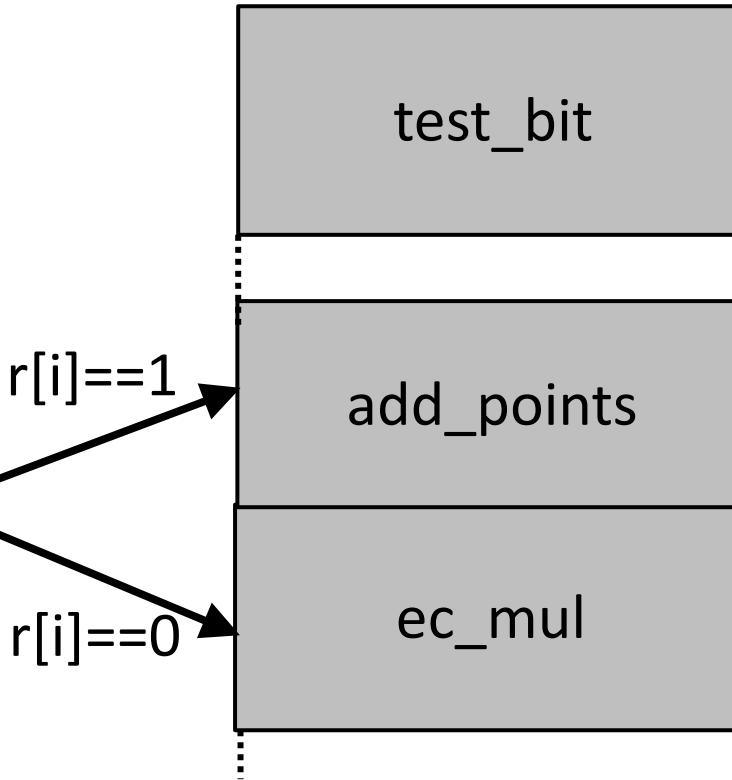


OS can extract bits of private AES key

# Attack Example II: ECDSA Signing

ECDSA Signing Routine

```
ec_mul(r, G) {
    res = 0
    nbits = |r|
    for (i = nbits-1; i>=0; i--):
        res = dup_point(res)
    if (test_bit(r[i])):
        res = add_points(res, G)
    return res }
```



P3:  
0x9EB30

P2:  
0xA6CB0

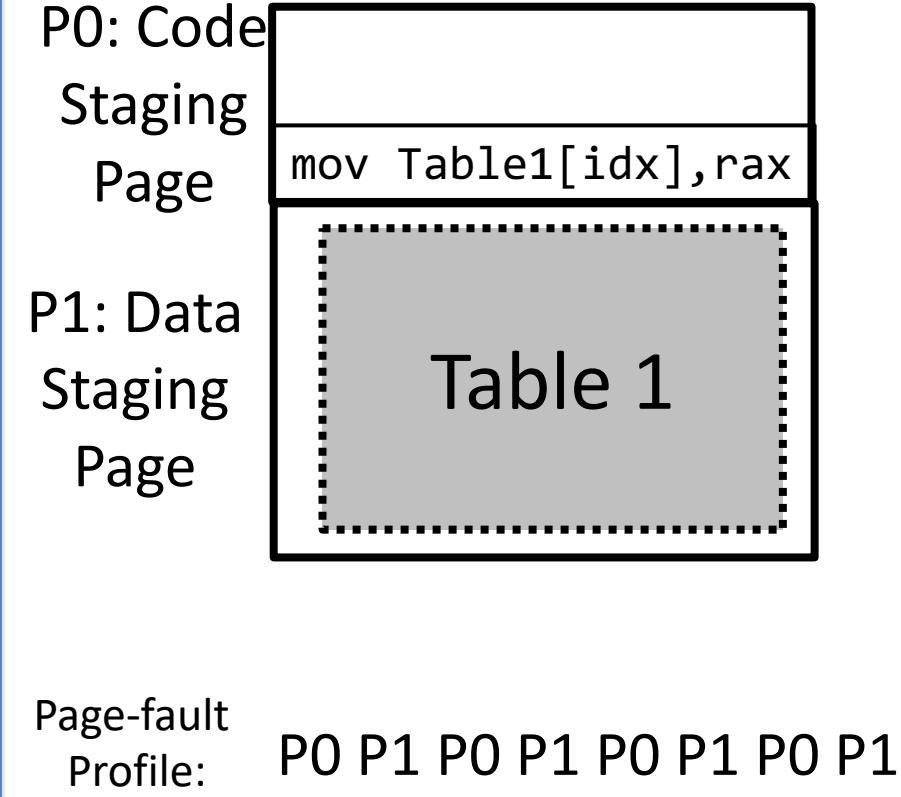
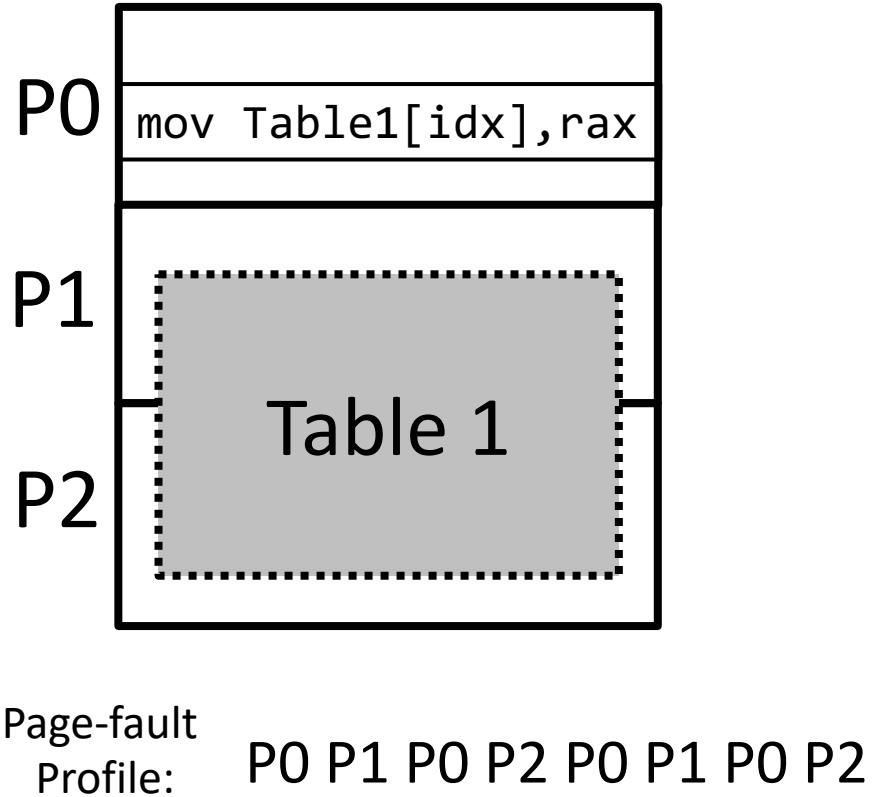
P1:  
0xA7310

Input dependent code page access reveals  
all the bits of ECDSA keys!

# Defenses against the PF side-channel

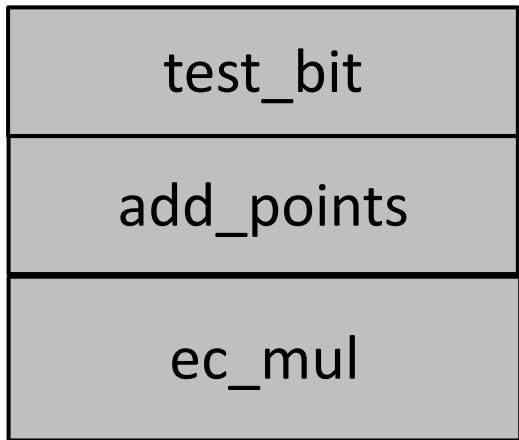
- Remove secret-dependent page-faults
  - Fit the table/branches in the same page
  - Always leads to a page fault
  - May not fit in one page
- Make the page fault pattern deterministic
  - Always leads to the same sequence of page faults
- Detect the attack in real-time and stop the leakage

# Defense Example: Determinising Data Accesses



# Example: Determinising Code Accesses

P3:  
0x9EB30  
  
P2:  
0xA6CB0  
  
P1:  
0xA7310



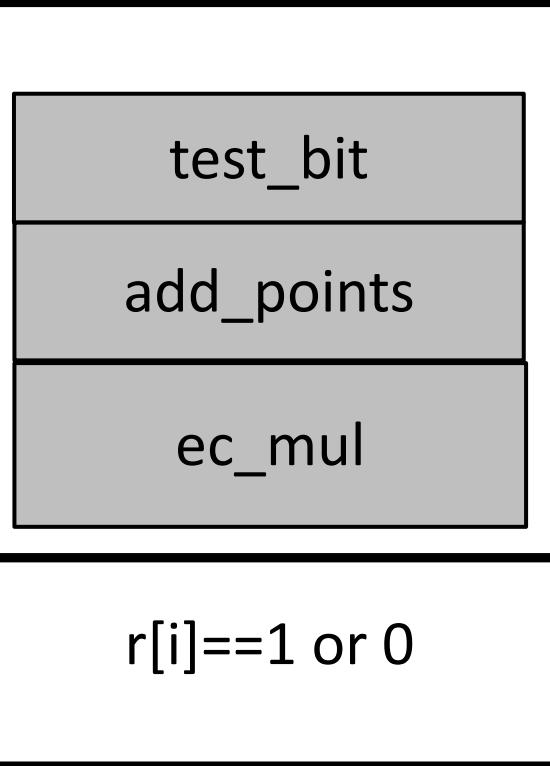
Page-fault  
Profile r=1  
  
Page-fault  
Profile r=0

P1 P2 P1 P3 P1 P1 P2  
  
P1 P2 P1 P3 P1

P0: Code  
Staging  
Page

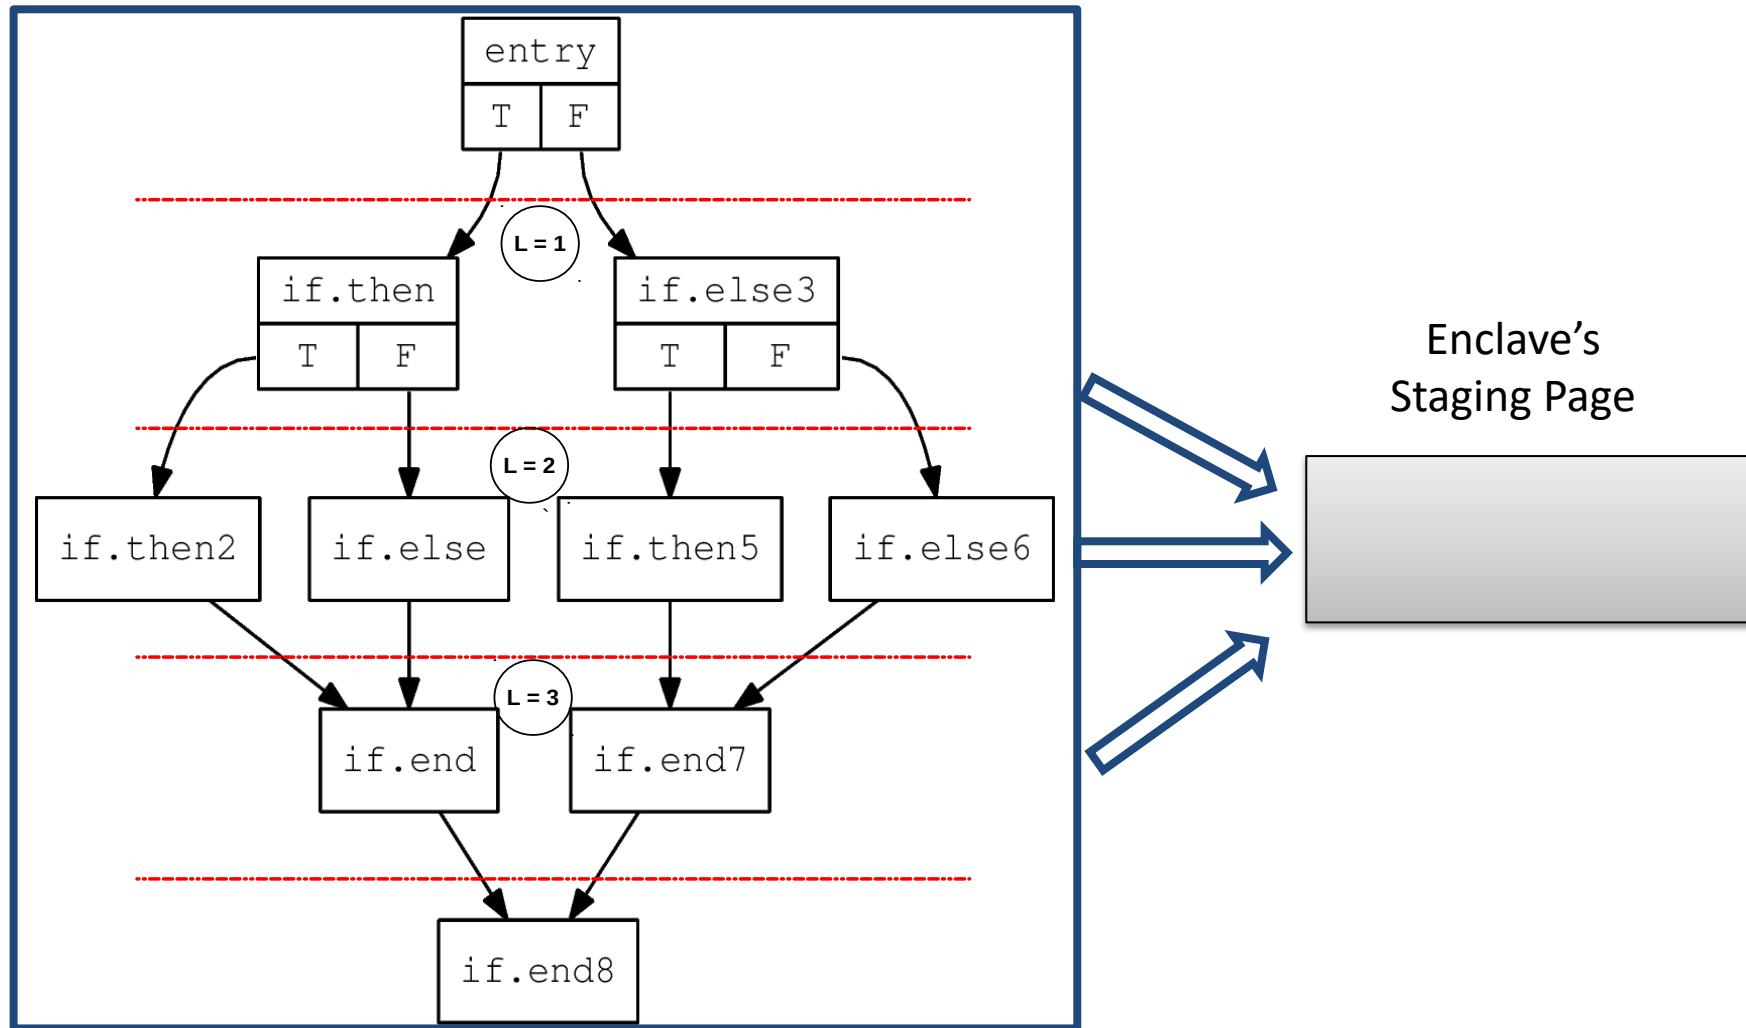
P1: Data  
Staging  
Page

Page-fault  
Profile:



P0 P1 P0

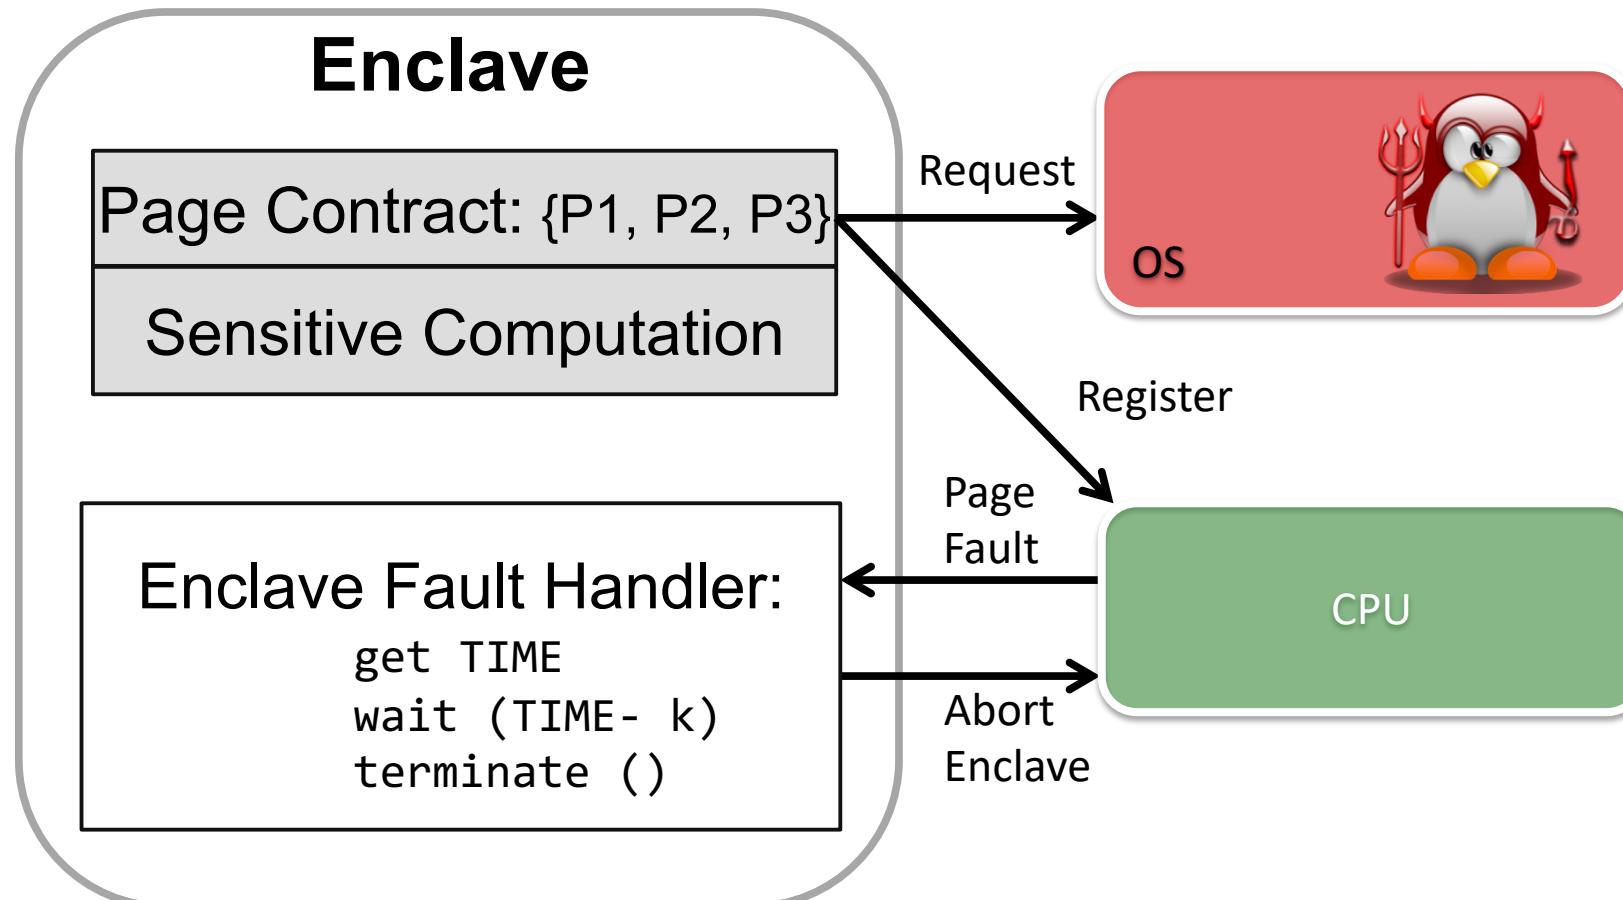
# Determinising Page Layout (multi-page)



# Multiplexed Execution

- Fetch all the input-dependent code / data to a staging page
- Execute the dependent-logic within the staging page via a multiplexer
- Save the updated values from data page

# Hardware Approach: Contractual Execution



# Summary

- 4 Primitives of Trusted Computing
  - Remote Attestation
  - Trusted Boot
  - Sealed Storage
  - Isolated Execution
- Using these primitives to understand Intel's TEE
- Identifying and fixing gaps in TEEs

# Shameless Plugs

- Join Secure & Trustworthy Systems (SECTRIS) group
  - Open positions for PhD & Master students.
  - Email me for more info
- Next Semester: 252-2603-00L Seminar on Systems Security

Thank you!

# Questions

- About the material covered in this lecture
- About the lab
  - Now is a good time to ask
  - More Q&A in the exercise and lab sessions
  - Discuss on Moodle